

# CSE7850/CX4803 Machine Learning in Computational Biology

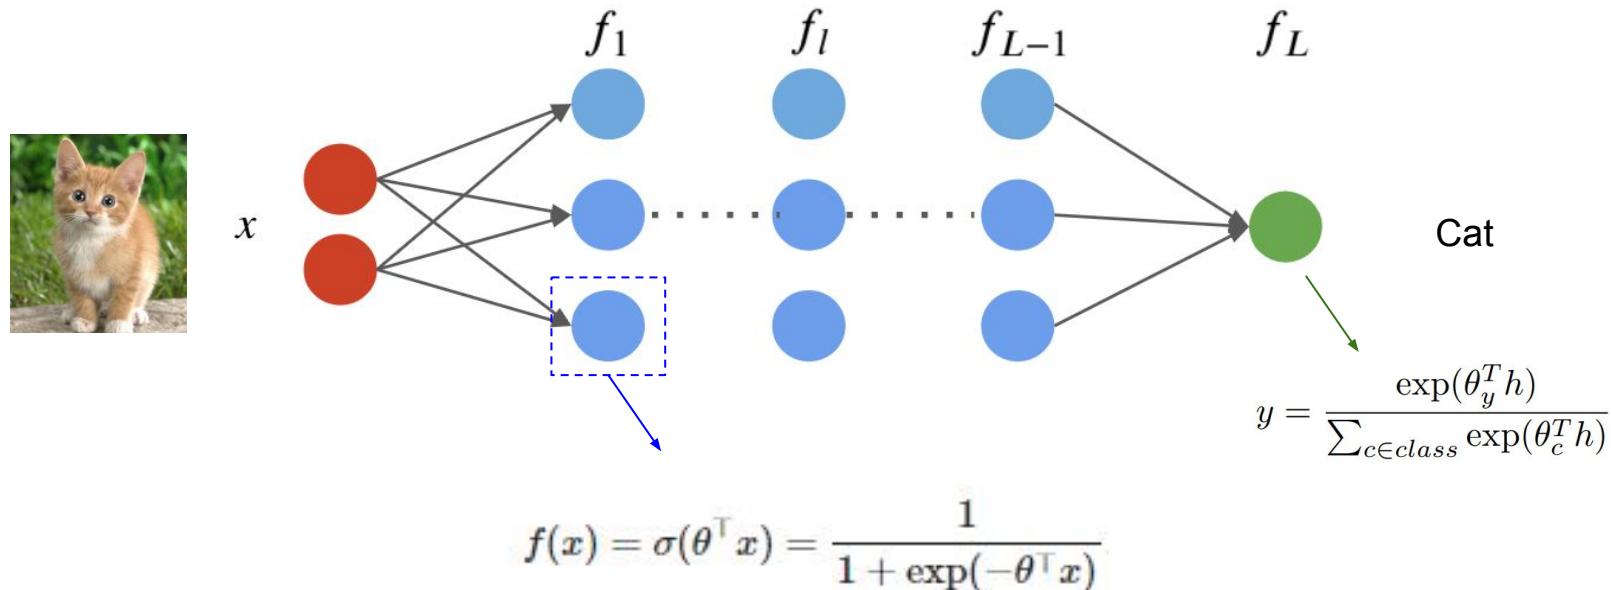


## Lecture 8: Deep Learning

Yunan Luo

# Recap: neural network

A (fully connected) neural network is a computational model that consists of a composition of multiple neural network layers



# Deep Learning

A class of machine learning methods that emphasizes:

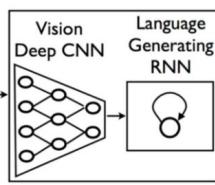
- Deep neural networks
- Large unstructured datasets, especially images, text, and audio
- Modern computational resources, like GPUs

# Expressivity of Deep Models

Deep neural networks are very powerful models in part because they can represent well complex mappings between  $x$  and  $y$ .

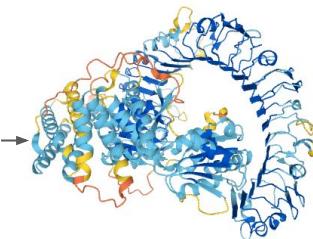
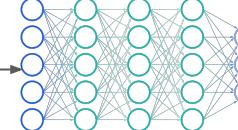
- Shallow neural networks can represent any function, but need very large hidden layers.
- Deep networks can represent very complex  $\mathcal{X} \rightarrow \mathcal{Y}$  mappings with fewer parameters.

In practice, deep neural networks can learn very complex mappings such as  $\text{image} \rightarrow \text{text description}$  that other algorithms cannot.



MAGELVSAFVNKLWDLLSHEYTLFQGVEDVAE  
CVEEIKDIVYDAEDVLETFVQKEKLGGTSGIRK  
RVIRDMQSFGVQQMIVDDYMHPLRNREERIRRT  
NYQVVSITGMGGLGKTTLARQVFNHDWMTKKFD  
EEETKEEEKKILEMTEYTLQRELYQLEMSKSL  
LLTSRNEISIVAPTNKYFNFKPECLKTDDSWKL  
IEHCGLPLAIKVLLGMIAEKYTSHWRRRLSEN  
FEELPSYKHCFLYLAHFPEDYIEKVENLSYYW  
VRNNIVISERDVKTTSRFETCHLHDMMIREVCLLK  
LVYQPTTLHVEKDINNPKLRSLVVTLGSWNM  
SCIGKLILHRLYSLEYAEVTHIPYSLGNLKLLI  
LALPSLIERTKLELSNLVKLETLENFSTKNSS

AlphaFold 2



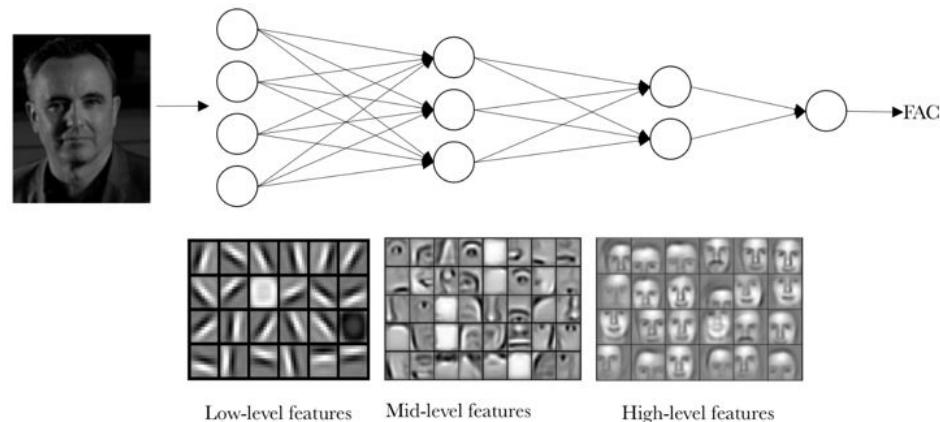
Ref: Learning CNN-LSTM Architectures for Image Caption Generation: <https://cs224d.stanford.edu/reports/msoh.pdf>

# Representation Learning

How does a deep neural network use its representational power?

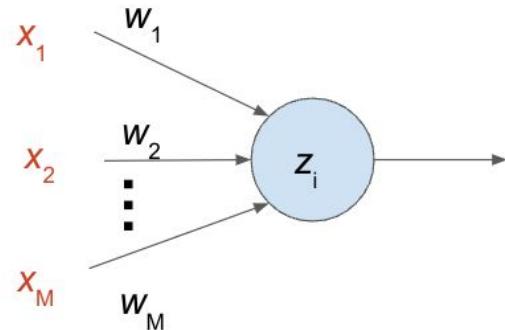
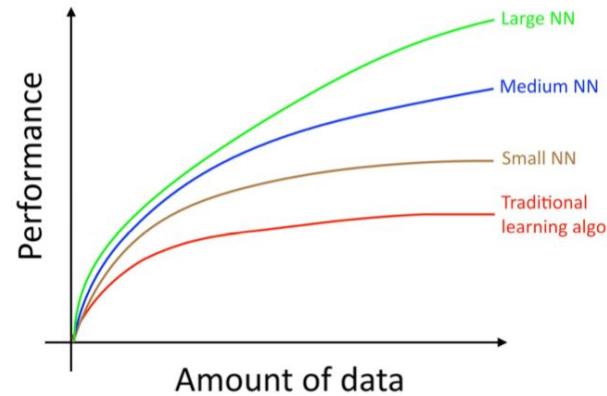
- The sequential layers can learn to represent data **at an increasing level of abstraction**
- This can also be interpreted as learning: each layer maps input to a more abstract feature space that is **learned** from data.

This is an example of representations that a deep neural network learns from data.



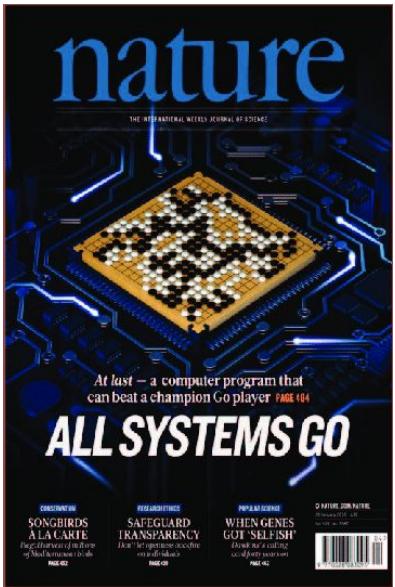
# Scaling to Large Datasets

- Deep neural nets also **scale to very large datasets**
  - Classical algorithms like linear regression saturate after a certain dataset size. Deep learning models typically keep improving as we add more data.
- Deep learning datasets benefit from large datasets because they can easily use specialized computational hardware.
  - Neural networks mostly implement **linear algebra**, and benefit from hardware acceleration.
  - Today, they are trained on graphical processing units (GPUs), or even more specialized hardware.

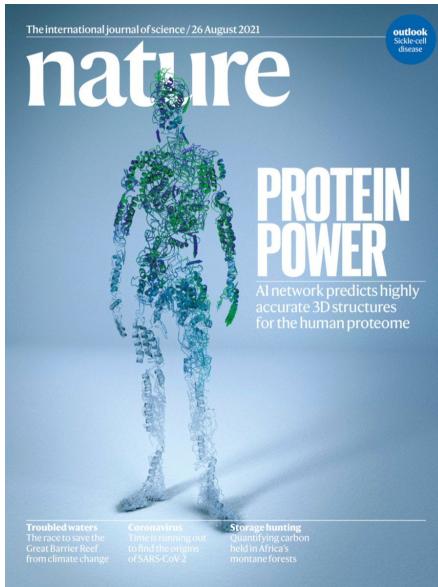


$$f_{\mathbf{w}}(\mathbf{x}) = \sigma \left( \sum_i w_i \cdot x_i \right) = \sigma(\mathbf{w}^T \mathbf{x})$$

# Successes of Deep Learning



Alpha Go (2016)



Alpha Fold (2018-2021)

A screenshot of the ChatGPT interface. The top right corner features the OpenAI logo. The main area contains the text 'ChatGPT: Optimizing Language Models for Dialogue' in large, bold, white font. Below this, a block of text in white describes the model's capabilities. The background is dark grey.

ChatGPT (2022)

# Successes of Deep Learning

Google Translate

The screenshot shows the Google Translate interface. The source language is set to "English - detected" and the target language is "Chinese (Simplified)". The input text "deep learning" is translated into "深度学习" (Shèndù xuéxí). Below the translation, there are two small icons: a microphone for voice input and a speaker for audio output. At the bottom of the interface, there are links for "Open in Google Translate" and "Feedback".

English - detected

Chinese (Simplified)

deep learning

深度学习  
Shèndù xuéxí

Open in Google Translate · Feedback

Translation



Speech recognition

# Successes of Deep Learning

A beautiful field of vibrant flowers  
atmospheric, hyper realistic, 8k, epic  
composition, cinematic, octane render,  
artstation landscape vista photography  
by Carr Clifton & Galen Rowell, 16K  
resolution, Landscape veduta photo by  
Dustin Lefevre & tdraw, 8k resolution,  
detailed landscape painting by Ivan  
Shishkin, DeviantArt, Flickr, rendered  
in Enscape, Miyazaki, Nausicaa Ghibli,  
Breath of The Wild, 4k detailed post  
processing, artstation, rendering by  
octane, unreal engine —ar 16:9



Text-to-image generative models  
(Stable diffusion, DALL-E, Imagen, ...)

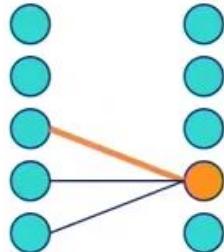


Image from: <https://prompthero.com/prompt/722b1199390>

# Pros and Cons of Deep Neural Networks

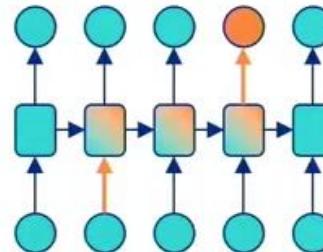
- **Pros:** Deep neural nets (DNNs) are among the most powerful ML models.
  - DNNs set a new state-of-the-art in many application areas.
  - They can be combined with many algorithms (supervised, unsupervised), and significantly improve them.
  - Many specialized software and hardware platforms for DNNs.
- **Cons:** A number of technical challenges:
  - DNNs can be slow and hard to train and require a lot of data
  - Deep neural networks are prone to overfitting.
  - Requires a massive amount of data

# Deep Learning Models and Implementations



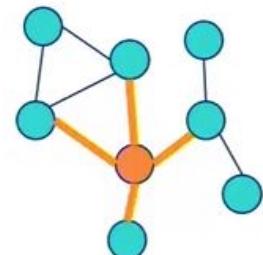
Convolutional Networks  
(e.g. computer vision)

- data in regular grid
- information flow to local neighbours
- AlphaFold 1



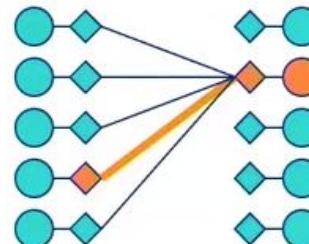
Recurrent Networks  
(e.g. language)

- data in ordered sequence
- information flow sequentially



Graph Networks (e.g. recommender systems or molecules)

- data in fixed graph structure
- information flow along fixed edges



Attention Module (e.g. language)

- data in unordered set
- information flow dynamically controlled by the network (via keys and queries)

PyTorch

TensorFlow



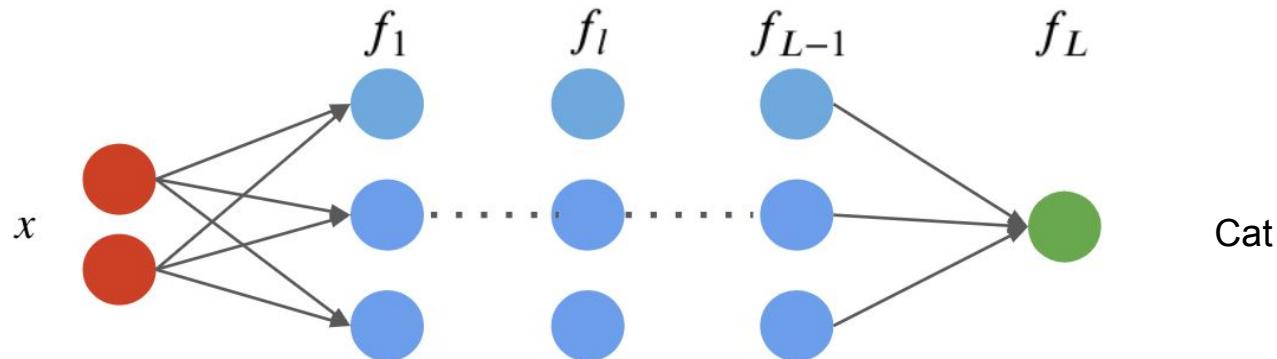
Keras

Google colab

An example of Deep Learning Models:  
**Convolutional Neural Network (CNN)**

# Background: Computer Vision

# Image Classification



# Image Classification

- ImageNet LSVRC-2011 contest:
  - **Dataset:** 1.2 million labeled images, 1000 classes
  - **Task:** Given a new image, label it with the correct class
  - **Multiclass** classification problem
- Examples from <http://image-net.org/>

Bird

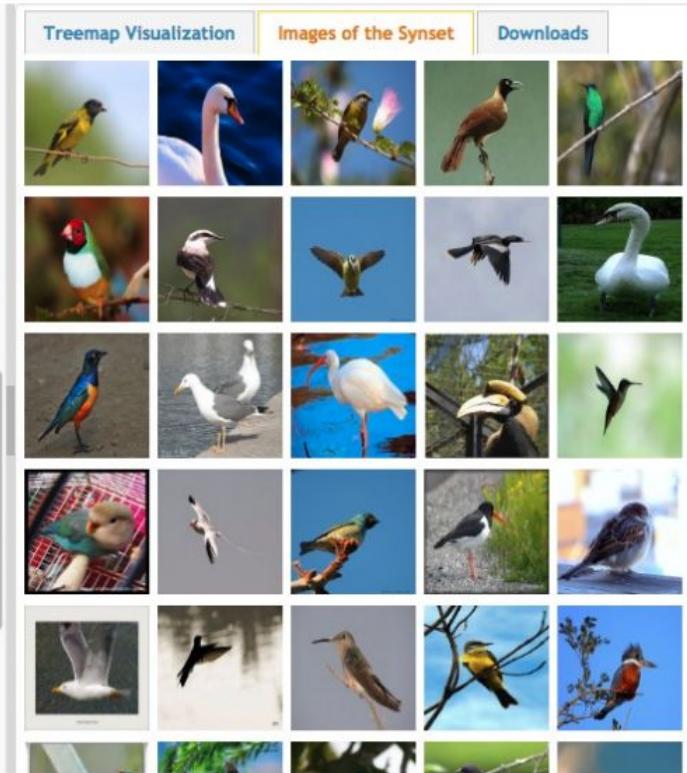
Warm-blooded egg-laying vertebrates characterized by feathers and forelimbs modified as wings

2126  
pictures

**92.85%**  
Popularity  
Percentile



- marine animal, marine creature, sea animal, sea creature (1)
  - scavenger (1)
  - biped (0)
  - predator, predatory animal (1)
  - larva (49)
  - acrodont (0)
  - feeder (0)
  - stunt (0)
  - chordate (3087)
    - tunicate, urochordate, urochord (6)
    - cephalochordate (1)
    - vertebrate, craniate (3077)
      - mammal, mammalian (1169)
      - bird (871)
        - dickeybird, dickey-bird, dickybird, dicky-bird (0)
        - cock (1)
        - hen (0)
        - nester (0)
        - night bird (1)
        - bird of passage (0)
        - protoavis (0)
        - archaeopteryx, archeopteryx, Archaeopteryx lithographica (0)
        - Sinornis (0)
        - Ibero-mesornis (0)
        - archaeornis (0)
        - ratite, ratite bird, flightless bird (10)
        - carinate, carinate bird, flying bird (0)
        - passerine, passeriform bird (279)
        - nonpasserine bird (0)
        - bird of prey, raptor, raptorial bird (80)
        - gallinaceous bird, gallinacean (114)



## German iris, Iris kochii

Iris of northern Italy having deep blue-purple flowers; similar to but smaller than Iris germanica

469 pictures

49.6%  
Popularity  
Percentile Wordnet  
IDs

- halophyte (0)
- succulent (39)
- cultivar (0)
- cultivated plant (0)
- weed (54)
  - evergreen, evergreen plant (0)
  - deciduous plant (0)
- vine (272)
- creeper (0)
- woody plant, ligneous plant (1868)
- geophyte (0)
- desert plant, xerophyte, xerophytic plant, xerophile, xerophilic mesophyte, mesophytic plant (0)
- aquatic plant, water plant, hydrophyte, hydrophytic plant (11)
- tuberous plant (0)
- bulbous plant (179)
  - + iridaceous plant (27)
    - + iris, flag, fleur-de-lis, sword lily (19)
      - bearded iris (4)
        - Florentine iris, orris, Iris germanica florentina, Iris German iris, Iris germanica (0)
        - German iris, Iris kochii (0)
        - Dalmatian iris, Iris pallida (0)
      - beardless iris (4)
      - bulbous iris (0)
      - dwarf iris, Iris cristata (0)
      - stinking iris, gladdon, gladdon iris, stinking gladwyn, Persian iris, Iris persica (0)
      - yellow iris, yellow flag, yellow water flag, Iris pseudacorus (0)
      - dwarf iris, vernal iris, Iris verna (0)
      - blue flag, Iris versicolor (0)

Treemap Visualization



## Court, courtyard

An area wholly or partly surrounded by walls or buildings; "the house was built around an inner court"

165 pictures

92.61% Popularity Percentile

 Wordnet IDs

Numbers in brackets: (the number of synsets in the subtree ).

- ↳ ImageNet 2011 Fall Release (32326)
  - plant, flora, plant life (4486)
  - geological formation, formation (175)
  - natural object (1112)
  - sport, athletics (176)
  - ↳ artifact, artefact (10504)
    - instrumentality, instrumentation (5494)
    - ↳ structure, construction (1405)
      - airdock, hangar, repair shed (0)
      - altar (1)
      - arcade, colonnade (1)
      - arch (31)
      - ↳ area (344)
        - aisle (0)
        - auditorium (1)
        - baggage claim (0)
        - box (1)
        - breakfast area, breakfast nook (0)
        - bullpen (0)
        - chancel, sanctuary, bema (0)
        - choir (0)
        - corner, nook (2)
        - ↳ court, courtyard (6)
          - atrium (0)
          - bailey (0)
          - cloister (0)
          - food court (0)
          - forecourt (0)
          - parvis (0)

Treemap Visualization Images of the Synset Downloads

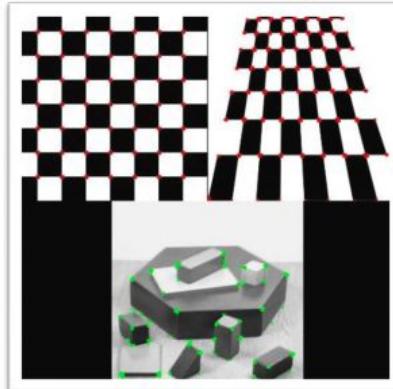


# Feature Engineering for Image Classification

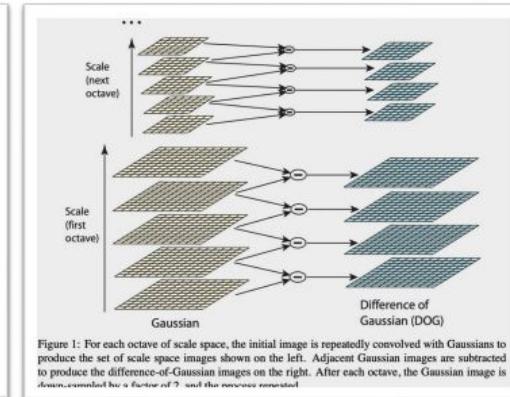
Edge detection (Canny)



Corner Detection (Harris)



Scale Invariant Feature Transform (SIFT)



Figures from <http://opencv.org>

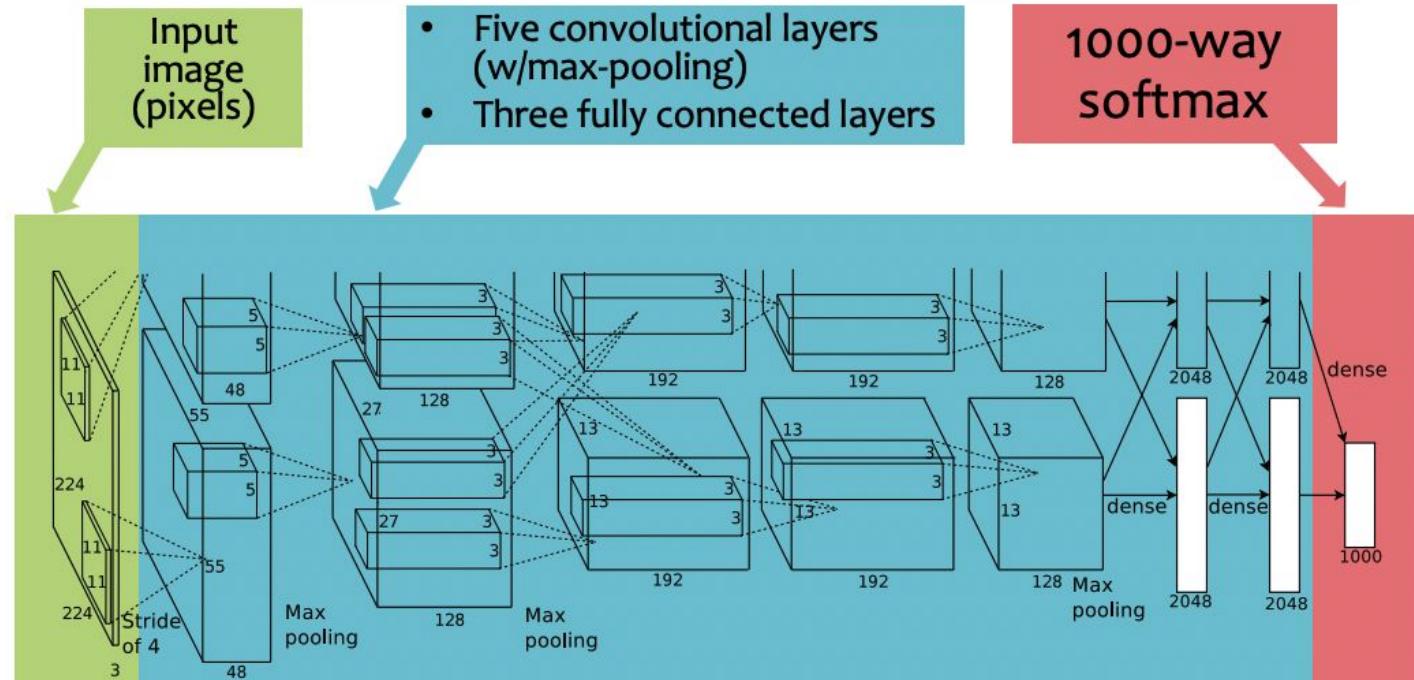
Figure from Lowe (1999) and Lowe (2004)

# Deep Learning for Image Classification (2012)

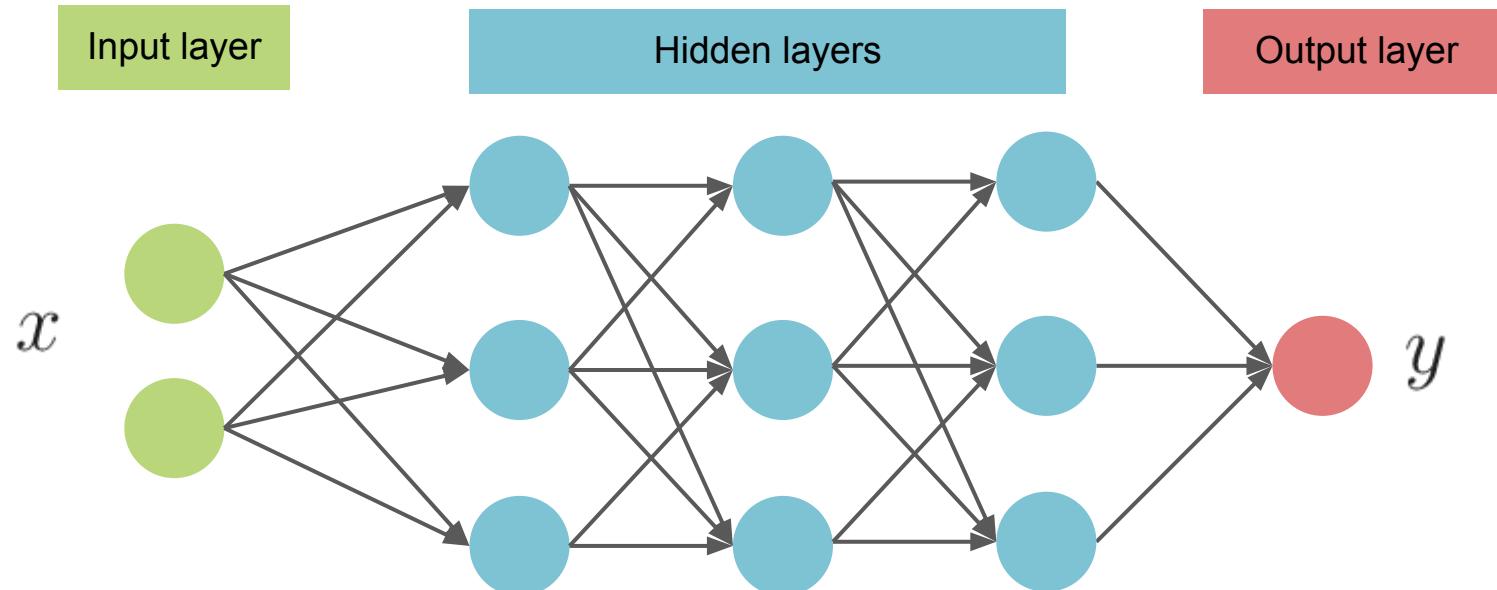
## CNN for Image Classification

(Krizhevsky, Sutskever & Hinton, 2012)

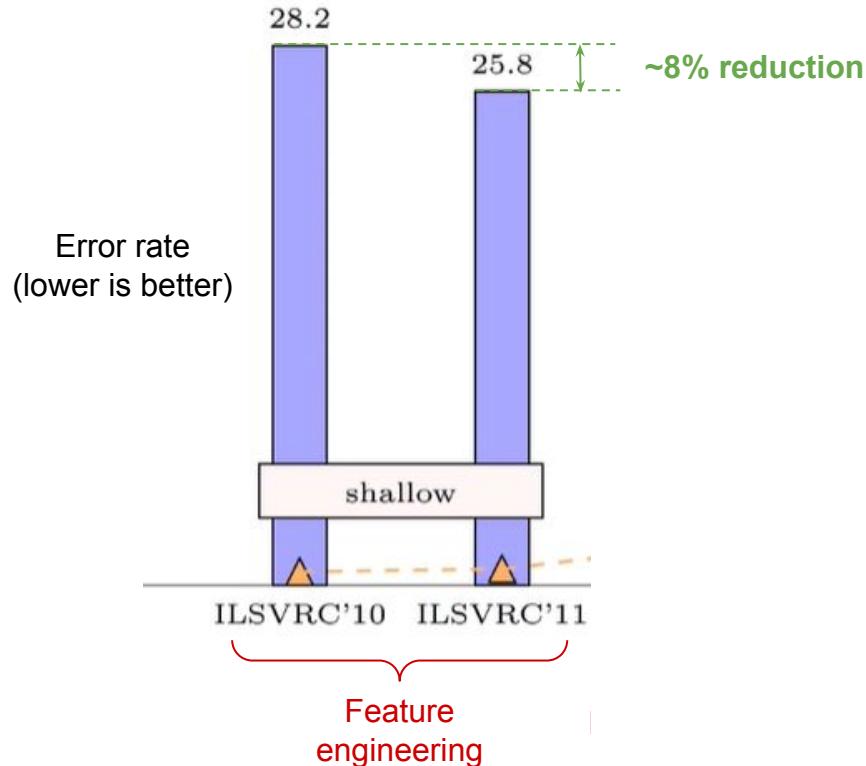
15.3% error on ImageNet LSVRC-2012 contest



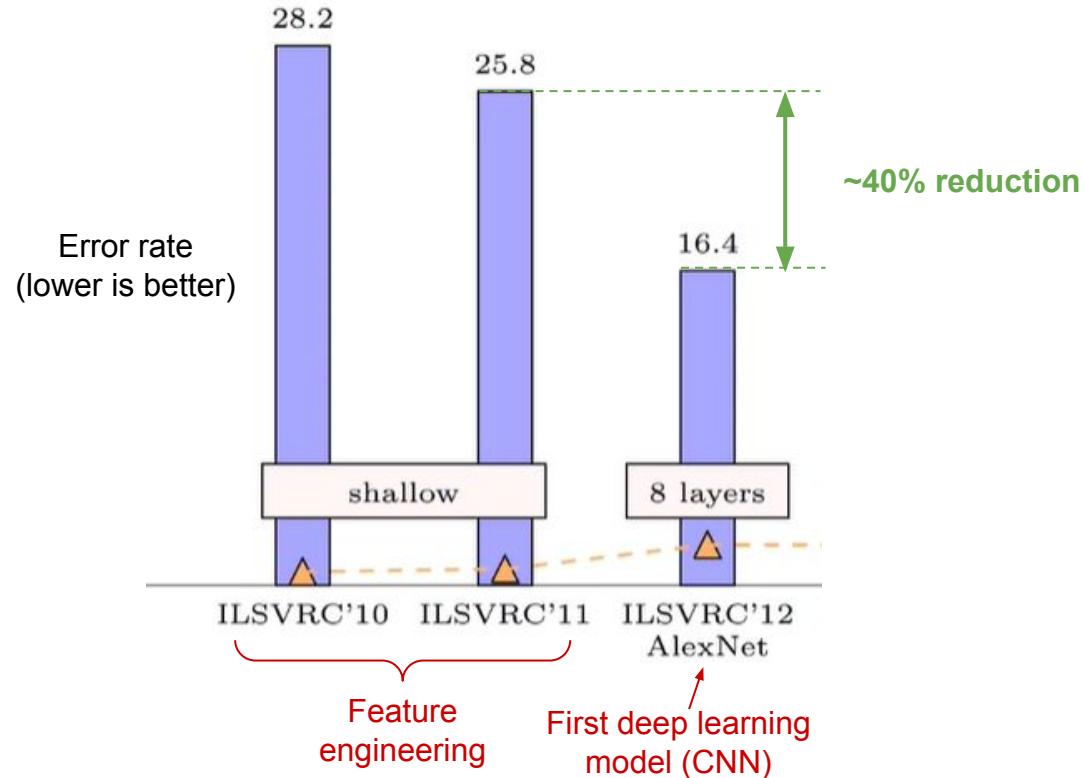
# Same idea!



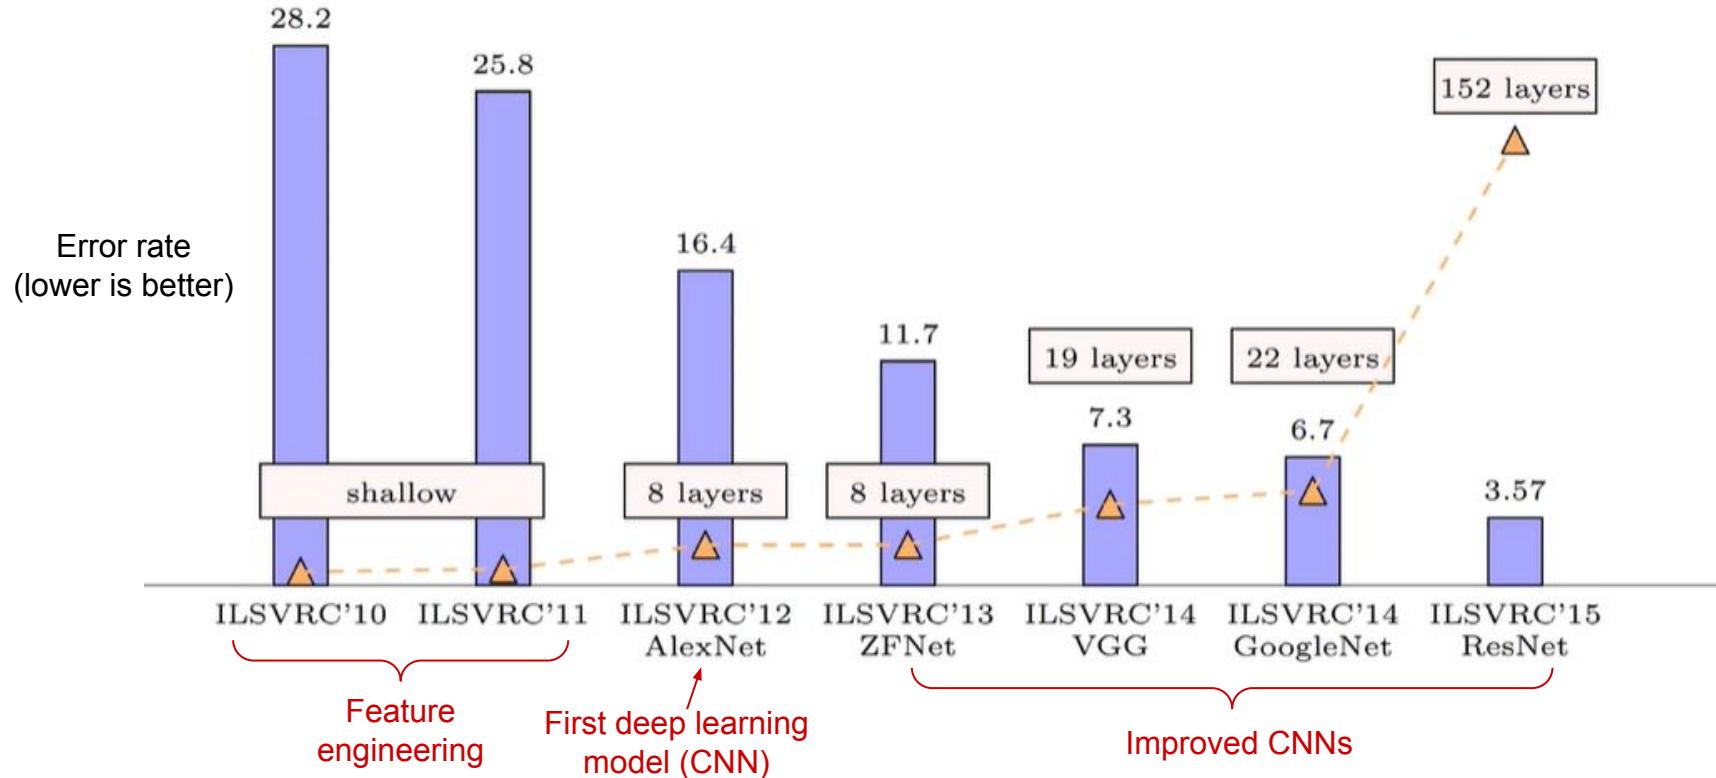
# How Deep Learning Performed on the ImageNet Contest?



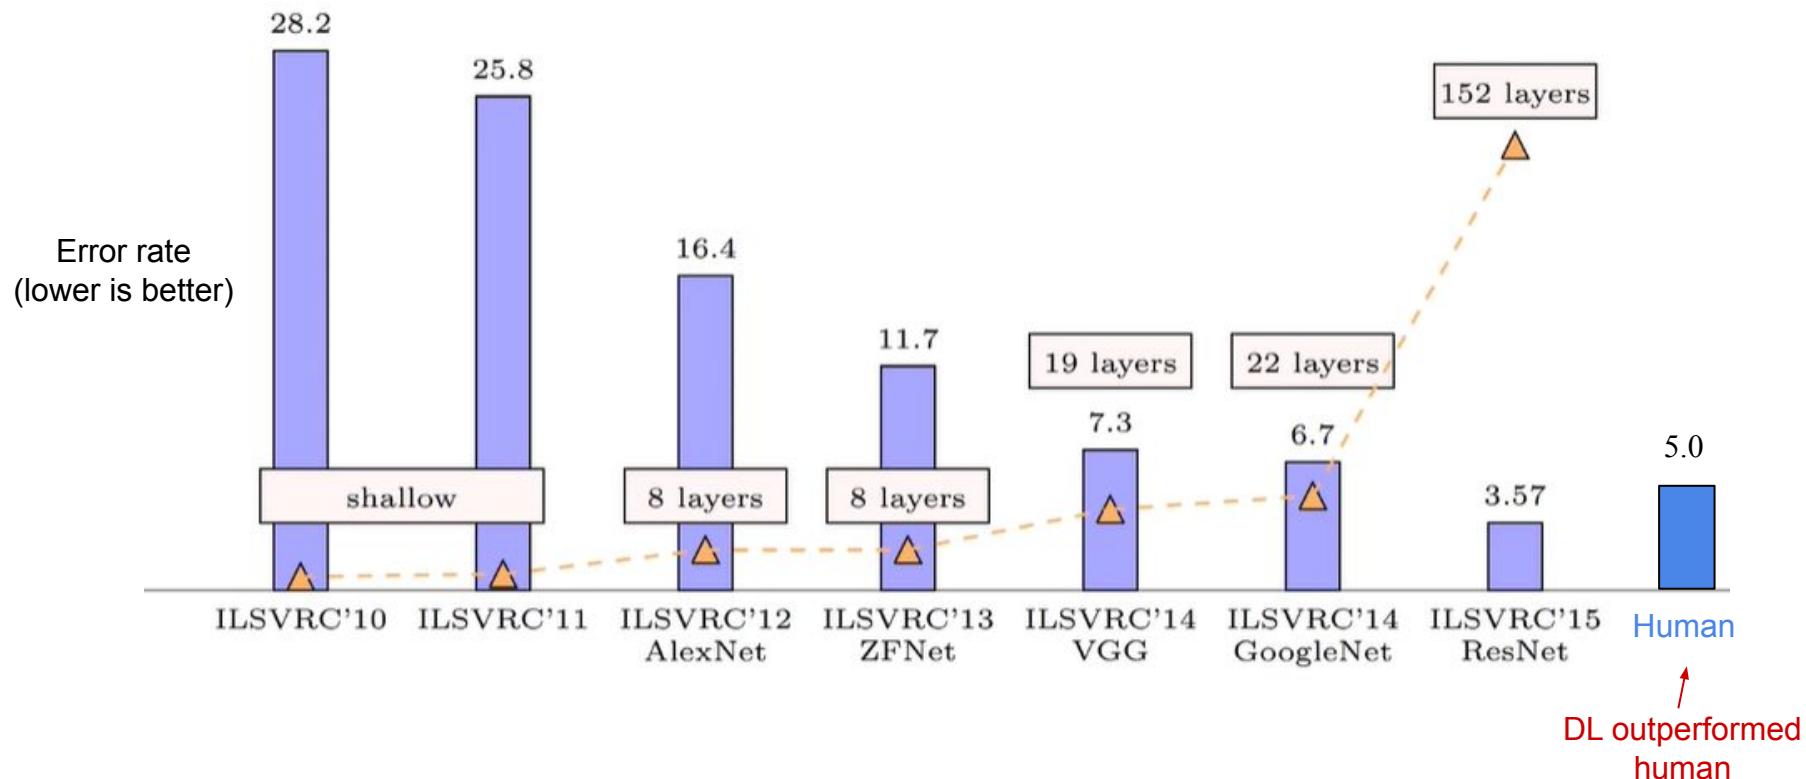
# How Deep Learning Performed on the ImageNet Contest?



# How Deep Learning Performed on the ImageNet Contest?



# How Deep Learning Performed on the ImageNet Contest?

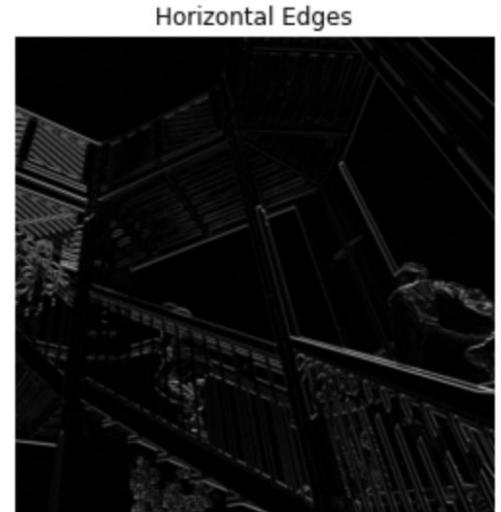
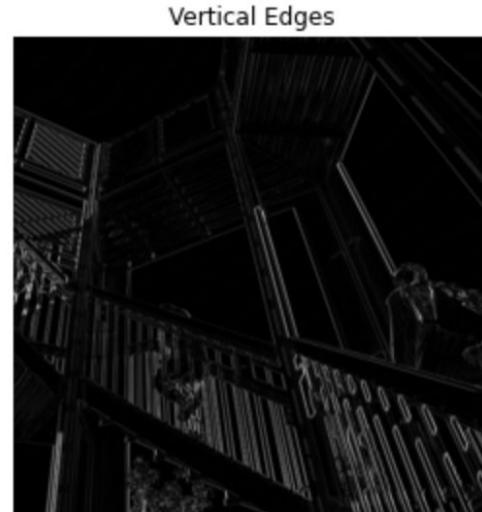


# Convolutional Neural Networks

- Our goal will be to introduce a modern and very powerful deep neural network architecture: the convolutional neural network (CNN).
  - CNNs are state-of-the-art across image processing tasks
  - They are loosely motivated by the structure of the visual cortex
  - They are based on two operations: **convolution** and **pooling**

# Motivation: Edge Detection

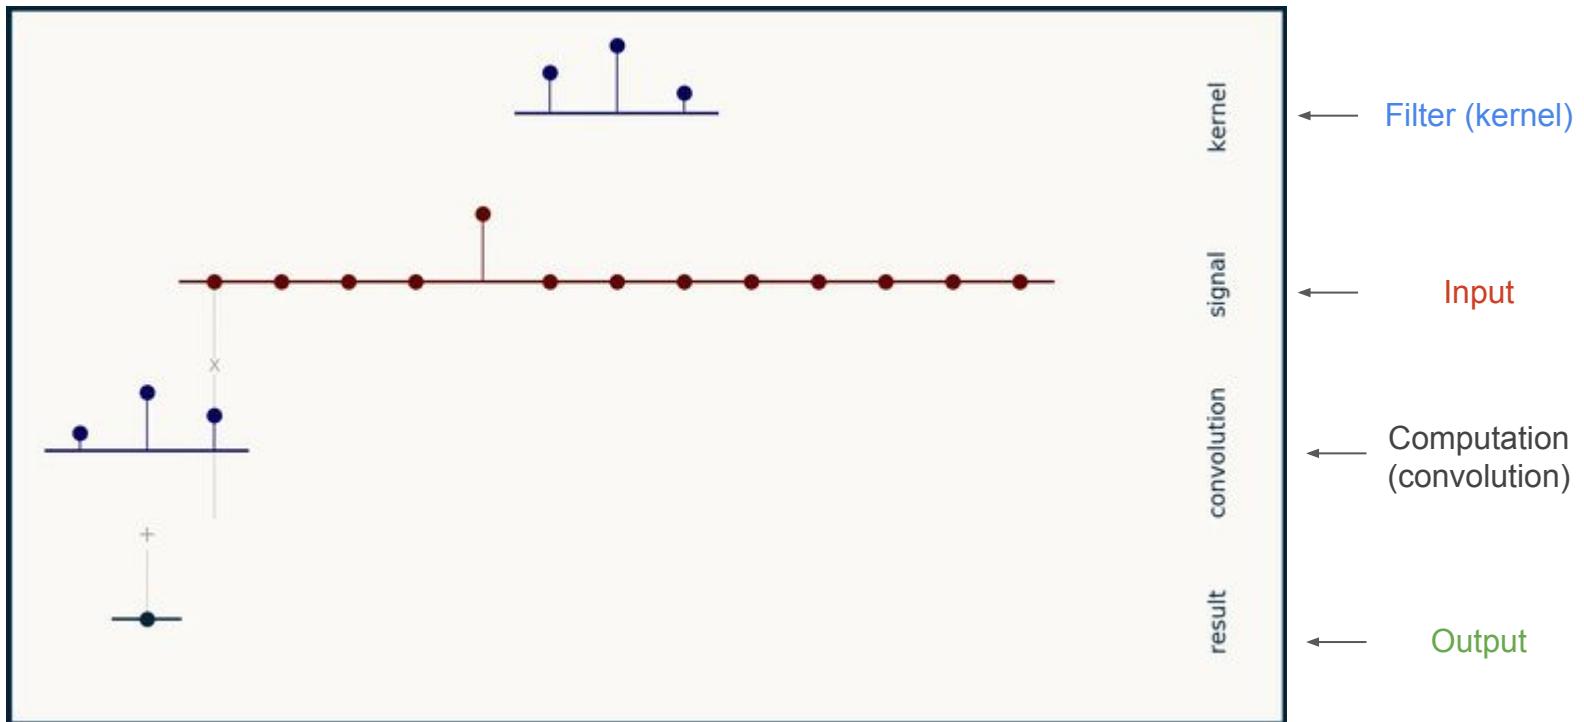
- Consider the problem of detecting edges in an image.



- This can be useful for detecting or classifying objects in a scene.

# Convolution: Intuition

- Here, we convolve a  $1 \times 3$  *filter* with a  $1 \times 13$  *signal* to obtain an *activation map*.



- Here, we convolve a  $1 \times 3$  filter with a  $1 \times 13$  signal to obtain an activation map.

# Convolution: definition

Let  $f \in \mathbb{R}^p$  and  $x \in \mathbb{R}^d$  be two vectors, called the *filter* (or *kernel*) and the *signal* respectively. Typically,  $p < d$ .

In deep learning, a convolution  $(f * x) \in \mathbb{R}^d$  is typically defined as

$$(f * x)[i] \triangleq \underbrace{\sum_{j=1}^p f[j]x[i+j]}_{\text{dot product of } f \text{ with part of } g} \quad (1)$$

where  $x[j] = 0$  when  $j \leq 0$  or  $j > d$ .

- We may think of the convolution

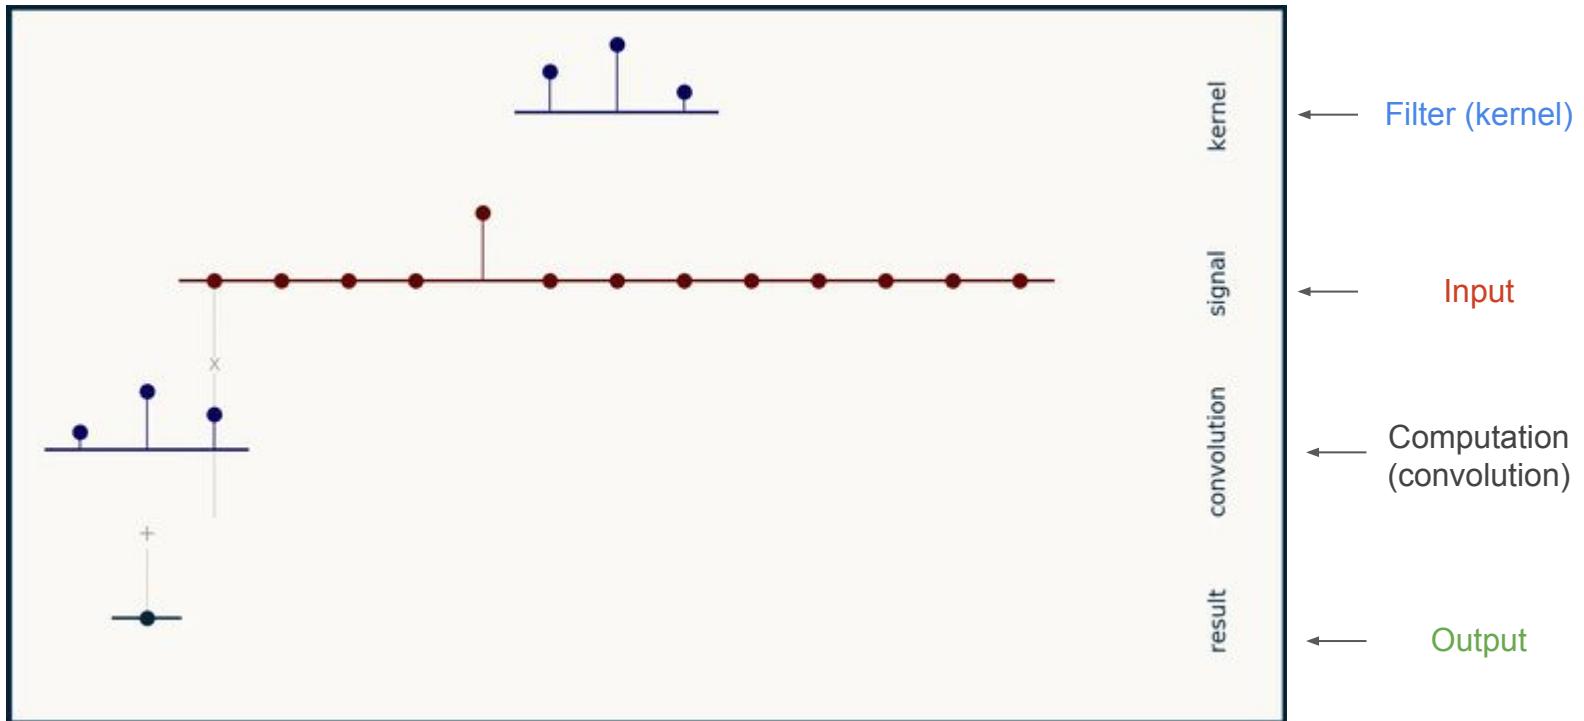
$$\sum_{j=1}^p f[j]x[i+j]$$

as a "local" dot product of `f` with `x[i:i+p]` (the subsequence of  $x$  that starts at position  $i$ ).

- $(f * x)[i]$  is large when  $f$  is similar to  $x$  at  $i$  and small otherwise. It measures the extent to which  $f$  is "found" in  $x$  at  $i$ .

# Convolution: Intuition

- Here, we convolve a  $1 \times 3$  *filter* with a  $1 \times 13$  *signal* to obtain an *activation map*.



- Notice how the output has large values at locations where the dot product between the kernel and the signal is large.

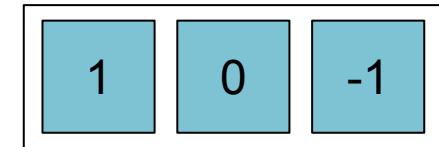
$$(f * x)[i] \triangleq \sum_{j=1}^p f[j]x[i+j]$$

# An example of 1D convolution

Input:



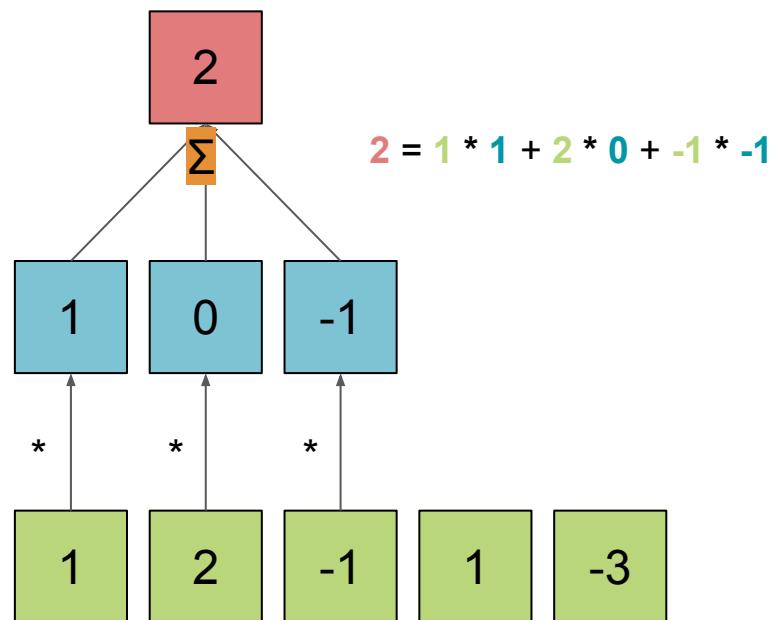
Filter:



$$(f * x)[i] \triangleq \sum_{j=1}^p f[j]x[i+j]$$

# An example of 1D convolution

Output:

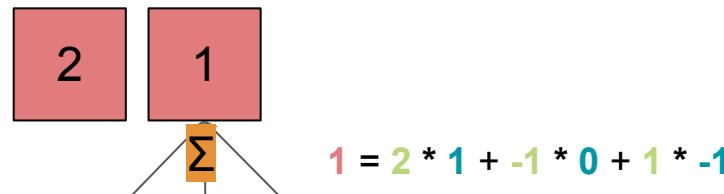


Filter:

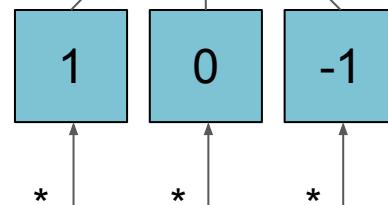
$$(f * x)[i] \triangleq \sum_{j=1}^p f[j]x[i+j]$$

# An example of 1D convolution

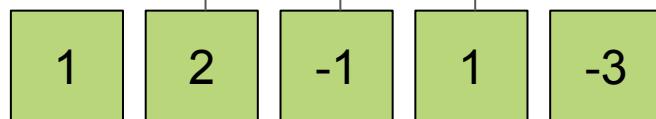
**Output:**



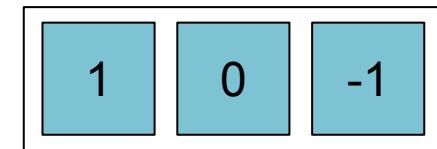
**Convolution:**



**Input:**



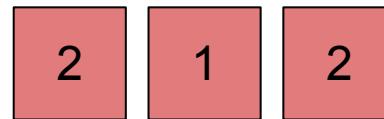
**Filter:**



$$(f * x)[i] \triangleq \sum_{j=1}^p f[j]x[i+j]$$

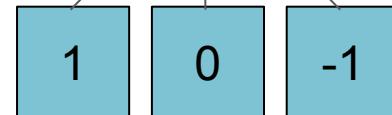
# An example of 1D convolution

Output:

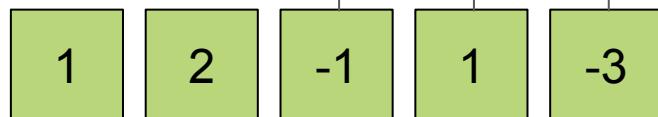


$$2 = -1 * 1 + 1 * 0 + -3 * -1$$

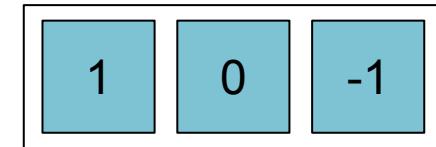
Convolution:



Input:

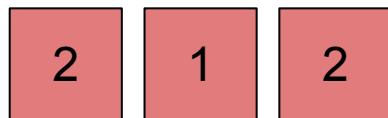


Filter:



# An example of 1D convolution

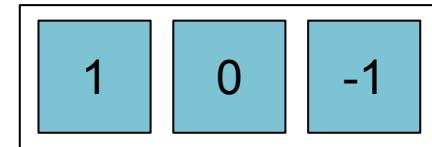
**Output:**



**Input:**



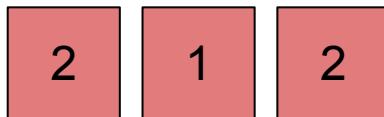
**Filter:**



- Different lengths (dimensions)
- We often want to keep the input and output of the CNN layer with the same dimension
- Use padding!

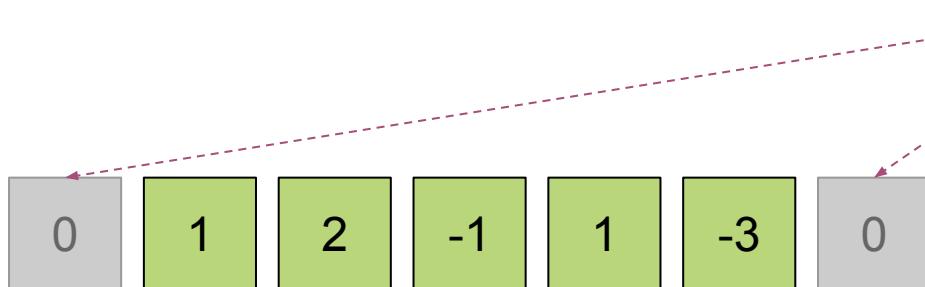
# An example of 1D convolution

**Output:**

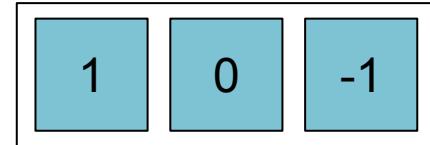


- Different lengths (dimensions)
- We often want to keep in the input and output of the CNN layer with the same dimension
- **Use padding!**

**Input:**

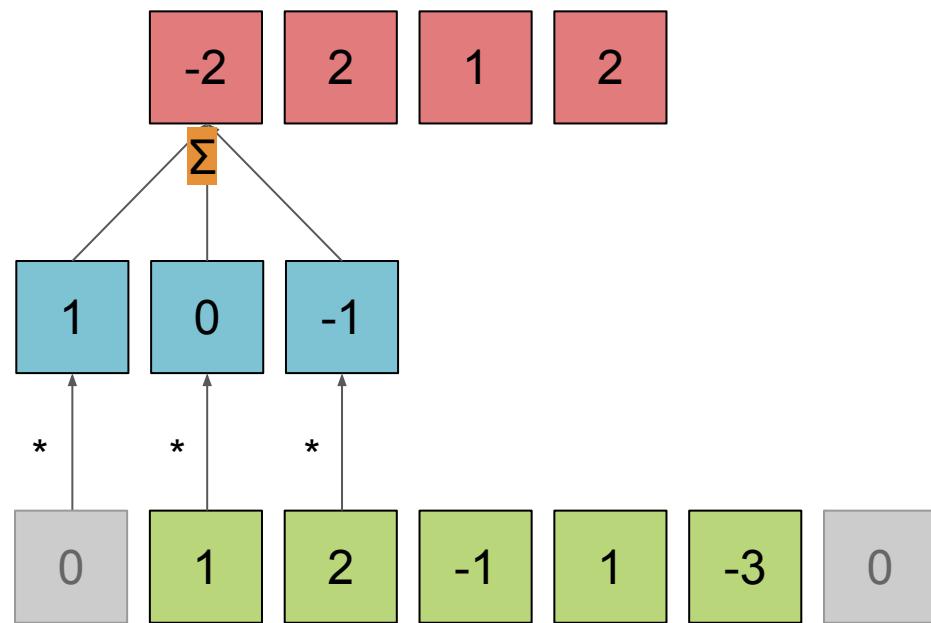


**Filter:**

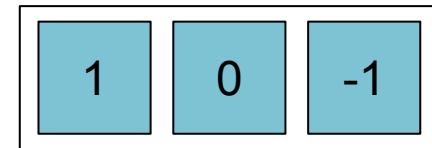


# An example of 1D convolution

Output:

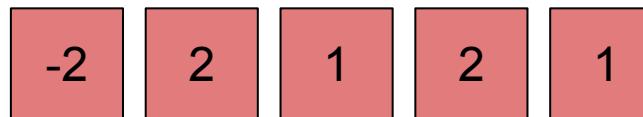


Filter:

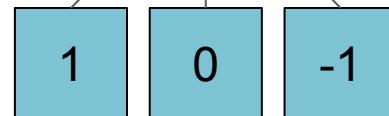


# An example of 1D convolution

Output:



$\Sigma$



Input:

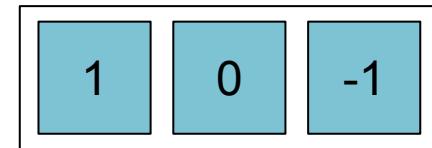


\*

\*

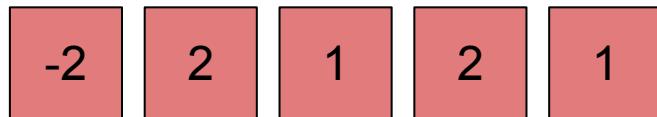
\*

Filter:

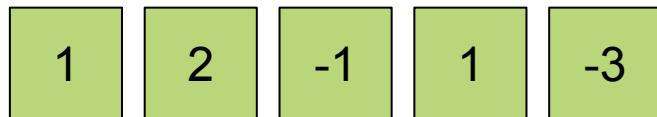


# An example of 1D convolution

**Output:**

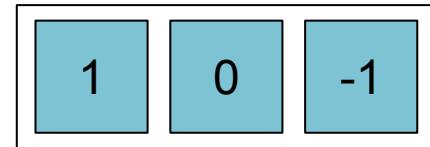


**Input:**

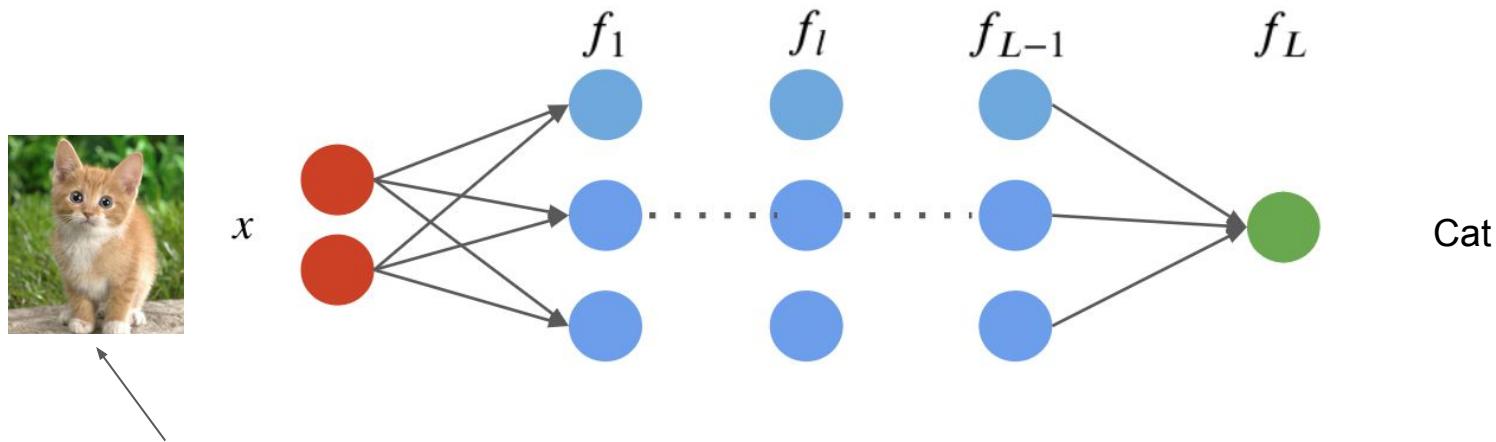


Same length (dimension)

**Filter:**



# Convolution on 2D data



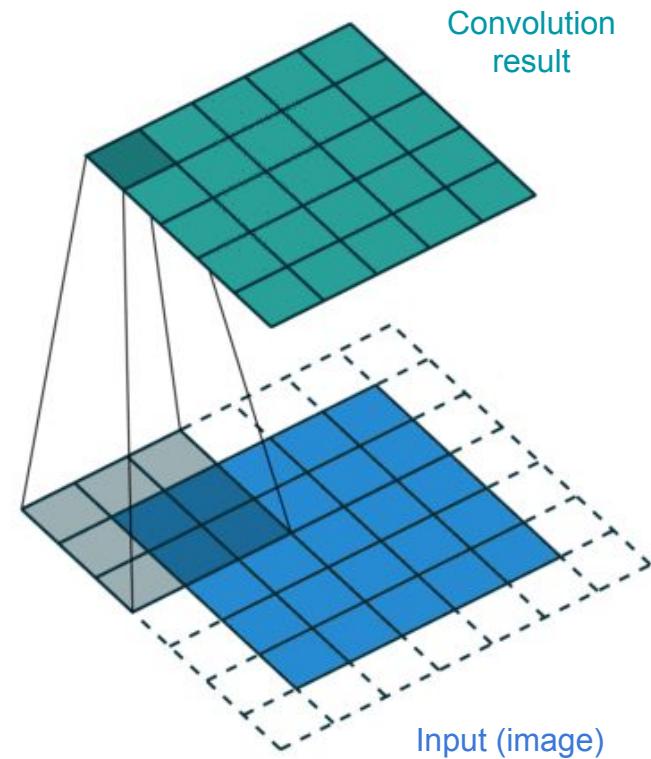
Need to represent an  
image in two dimensions

# 2D Convolutional Layers

In this example,  $f$  is a  $3 \times 3$  array (gray area).  $x$  is a  $5 \times 5$  array (blue area). The result of the convolution is the top green area.

In the context of neural networks,  $f$  are learnable weights,  $x$  are inputs from previous layers.

- The signal  $x$  can be an image.
- The filter  $f$  can be a "feature" of the image, like an edge or a more complex object.



# 2D Convolutional Layers

Input image

1	1	1	1	1
1	0	0	1	0
1	0	1	0	0
1	1	0	0	0
1	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

Output


# 2D Convolutional Layers

Input image

0	0	0	0	0	0	0
0	1	1	1	1	1	0
0	1	0	0	1	0	0
0	1	0	1	0	0	0
0	1	1	0	0	0	0
0	1	0	0	0	0	0
0	0	0	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

Output


# 2D Convolutional Layers

Input image						
0	0	0	0	0	0	0
0	1	1	1	1	1	0
0	1	0	0	1	0	0
0	1	0	1	0	0	0
0	1	1	0	0	0	0
0	1	0	0	0	0	0
0	0	0	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

Output


Compute “dot product”:

$$\begin{aligned} & 0 * 0 + 0 * 0 + 0 * 0 \\ & + 0 * 0 + \textcolor{green}{1} * \textcolor{blue}{1} + \textcolor{green}{1} * \textcolor{blue}{1} \\ & + 0 * 0 + \textcolor{green}{1} * \textcolor{blue}{1} + 0 * 0 \end{aligned}$$

# 2D Convolutional Layers

Input image								
0	0	0	0	0	0	0	0	0
0	1	1	1	1	1	1	0	0
0	1	0	0	1	0	0	0	0
0	1	0	1	0	0	0	0	0
0	1	1	0	0	0	0	0	0
0	1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

Output

3				

Compute “dot product”:

$$\begin{aligned} & 0 * 0 + 0 * 0 + 0 * 0 \\ & + 0 * 0 + \textcolor{green}{1} * \textcolor{blue}{1} + \textcolor{green}{1} * \textcolor{blue}{1} \\ & + 0 * 0 + \textcolor{green}{1} * \textcolor{blue}{1} + 0 * 0 \\ & = \textcolor{red}{3} \end{aligned}$$

# 2D Convolutional Layers

Input image							
0	0	0	0	0	0	0	0
0	1	1	1	1	1	1	0
0	1	0	0	1	0	0	0
0	1	0	1	0	0	0	0
0	1	1	0	0	0	0	0
0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

×

Output

3	2						

Compute “dot product”:

$$\begin{aligned} & 0 * 0 + 0 * 0 + 0 * 0 \\ & + 1 * 0 + 1 * 1 + 1 * 1 \\ & + 1 * 0 + 0 * 1 + 0 * 0 \\ & = 2 \end{aligned}$$

# 2D Convolutional Layers

Input image						
0	0	0	0	0	0	0
0	1	1	1	1	1	0
0	1	0	0	1	0	0
0	1	0	1	0	0	0
0	1	1	0	0	0	0
0	1	0	0	0	0	0
0	0	0	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

×

Output

3	2	2		

Compute “dot product”:

$$\begin{aligned} & 0 * 0 + 0 * 0 + 0 * 0 \\ & + 1 * 0 + 1 * 1 + 1 * 1 \\ & + 0 * 0 + 0 * 1 + 1 * 0 \\ & = 2 \end{aligned}$$

# 2D Convolutional Layers

Input image

0	0	0	0	0	0	0
0	1	1	1	1	1	0
0	1	0	0	1	0	0
0	1	0	1	0	0	0
0	1	1	0	0	0	0
0	1	0	0	0	0	0
0	0	0	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

Output

3	2	2	3	

# 2D Convolutional Layers

Input image							
0	0	0	0	0	0	0	0
0	1	1	1	1	1	1	0
0	1	0	0	1	0	0	0
0	1	0	1	0	0	0	0
0	1	1	0	0	0	0	0
0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

×

Output

3	2	2	3	1

# 2D Convolutional Layers

Input image

0	0	0	0	0	0	0	0
0	1	1	1	1	1	1	0
0	1	0	0	1	0	0	0
0	1	0	1	0	0	0	0
0	1	1	0	0	0	0	0
0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

Output

3	2	2	3	1
2				

# 2D Convolutional Layers

Input image							
0	0	0	0	0	0	0	0
0	1	1	1	1	1	0	
0	1	0	0	1	0	0	
0	1	0	1	0	0	0	
0	1	1	0	0	0	0	
0	1	0	0	0	0	0	
0	0	0	0	0	0	0	

Filter (kernel)

0	0	0
0	1	1
0	1	0

×

Output				
3	2	2	3	1
2	0			

# 2D Convolutional Layers

Input image						
0	0	0	0	0	0	0
0	1	1	1	1	1	0
0	1	0	0	1	0	0
0	1	0	1	0	0	0
0	1	1	0	0	0	0
0	1	0	0	0	0	0
0	0	0	0	0	0	0

Filter (kernel)

×

0	0	0
0	1	1
0	1	0

Output				
3	2	2	3	1
2	0	2	1	0
2	2	1	0	0
3	1	0	0	0
1	0	0	0	0

# 2D Convolutional Layers

Input image

1	1	1	1	1
1	0	0	1	0
1	0	1	0	0
1	1	0	0	0
1	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

Output

3	2	2	3	1
2	0	2	1	0
2	2	1	0	0
3	1	0	0	0
1	0	0	0	0

# 2D Convolutional Layers

Input image

0	0	0	0	0	0	0	0
0	1	1	1	1	1	1	0
0	1	0	0	1	0	0	0
0	1	0	1	0	0	0	0
0	1	1	0	0	0	0	0
0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

Output

3	2	2	3	1
2	0	2	1	0
2	2	1	0	0
3	1	0	0	0
1	0	0	0	0

The **output value is large** when  
the **local feature of the input**  
**match** the **feature in the filter**

# 2D Convolutional Layers

Input image							
0	0	0	0	0	0	0	0
0	1	1	1	1	1	1	0
0	1	0	0	1	0	0	0
0	1	0	1	0	0	0	0
0	1	1	0	0	0	0	0
0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

Output

3	2	2	3	1
2	0	2	1	0
2	2	1	0	0
3	1	0	0	0
1	0	0	0	0

The **output value is large** when  
the **local feature of the input**  
**match** the **feature in the filter**

# 2D Convolutional Layers

Input image						
0	0	0	0	0	0	0
0	1	1	1	1	1	0
0	1	0	0	1	0	0
0	1	0	1	0	0	0
0	1	1	0	0	0	0
0	1	0	0	0	0	0
0	0	0	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

Output

3	2	2	3	1
2	0	2	1	0
2	2	1	0	0
3	1	0	0	0
1	0	0	0	0

The **output value is large** when  
the **local feature of the input**  
**match** the **feature in the filter**

# 2D Convolutional Layers

Input image

1	1	1	1	1
1	0	0	1	0
1	0	1	0	0
1	1	0	0	0
1	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

We can use multiple filters to capture different patterns in the input

# 2D Convolutional Layers

Input image

1	1	1	1	1
1	0	0	1	0
1	0	1	0	0
1	1	0	0	0
1	0	0	0	0

Filter (kernel)

0	0	0
0	1	1
0	1	0

0	0	1
0	1	0
1	0	0

1	1	1
0	0	0
0	0	0

We can multiple filters to capture different patterns in the input

- In a CNN, we do not have to manually design those filters.
- The CNN will automatically learned those filters from data

# Pooling layer

- We have introduced convolutional layers. Next, we introduce another important component in a CNN -- **pooling layer**
- Example: Max Pooling

3	2	2	3
2	0	2	1
2	2	1	0
3	1	0	0

Max pool with 2x2 filters and stride 2



3	

# Pooling layer

- We have introduced convolutional layers. Next, we introduce another important component in a CNN -- **pooling layer**
- Example: Max Pooling

3	2	2	3
2	0	2	1
2	2	1	0
3	1	0	0

Max pool with 2x2 filters and stride 2



3	3

# Pooling layer

- We have introduced convolutional layers. Next, we introduce another important component in a CNN -- **pooling layer**
- Example: **Max** Pooling

3	2	2	3
2	0	2	1
2	2	1	0
3	1	0	0

Max pool with 2x2 filters and stride 2



3	3
3	1

# Pooling layer

- We have introduced convolutional layers. Next, we introduce another important component in a CNN -- **pooling layer**
- Example: Average Pooling

3	2	2	3
2	0	2	1
2	2	1	0
3	1	0	0

Average pool with 2x2 filters and stride 2



7/4	2
2	1/4

# Pooling layer

A pooling layer is a model  $f : \mathbb{R}^d \rightarrow \mathbb{R}^p$  that applies pooling operations to an input  $x$ .

$$f(x) = \text{pooling}(x)$$

where *pooling* is a pre-defined operation applied over the input  $x$ .

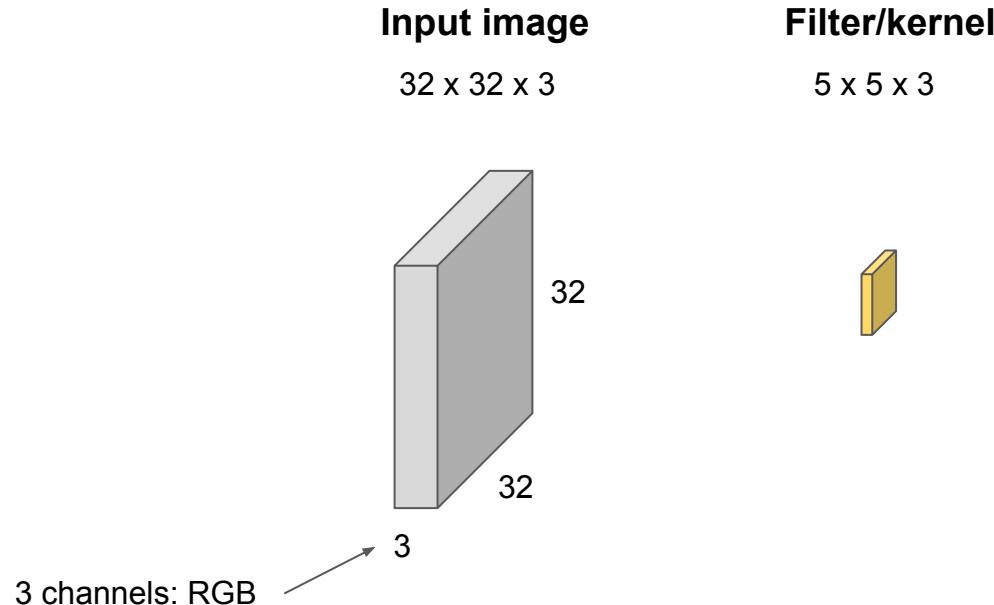
Pooling layer does not have learnable weights.

Purpose of Pooling:

- Allows slight **translation invariance** inside the pooling window: the exact location of the relevant feature is not critical. This improves feature detection.
- When used together with convolutions, pooling **reduces spatial size** of feature maps to capture more abstract representations.

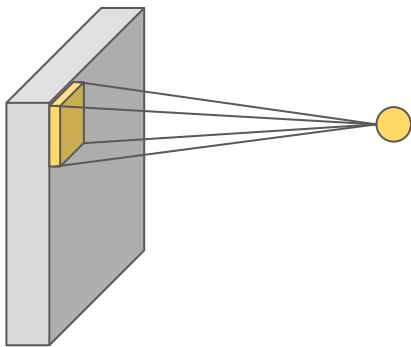
# Convolutional Neural Networks

# Applying Convolution on an RGB Image

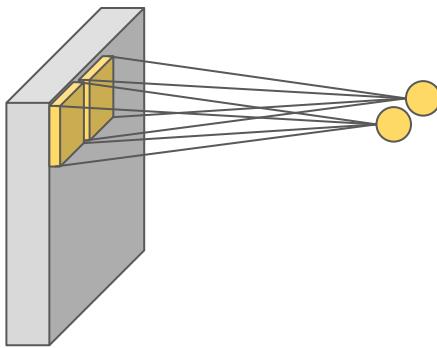


**Convolution:** slide the filter over the image spatially and compute the dot products

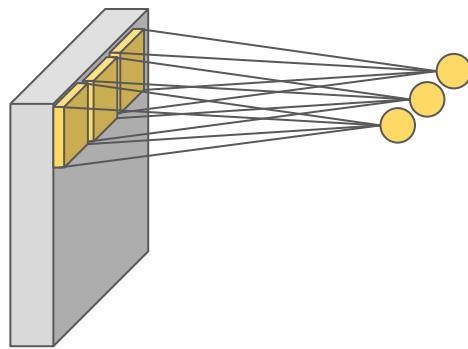
# Convolution layer



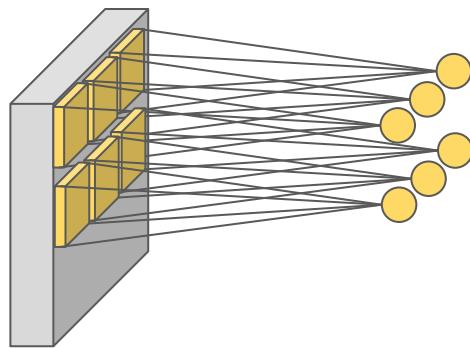
# Convolution layer



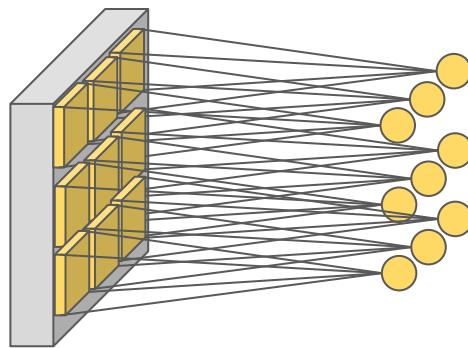
# Convolution layer



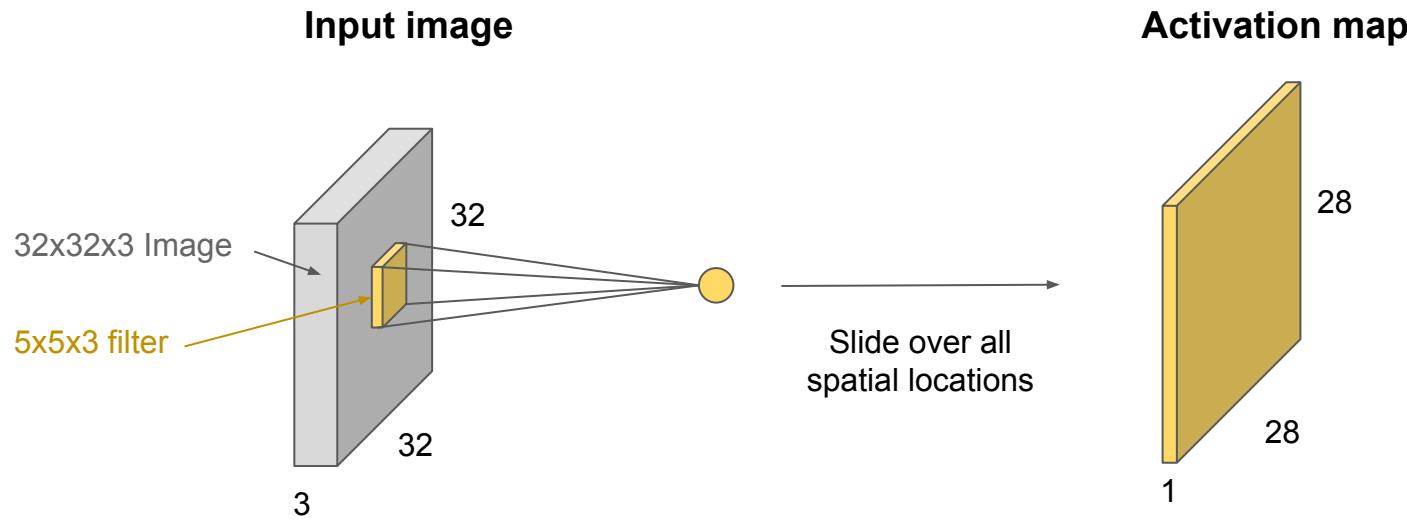
# Convolution layer



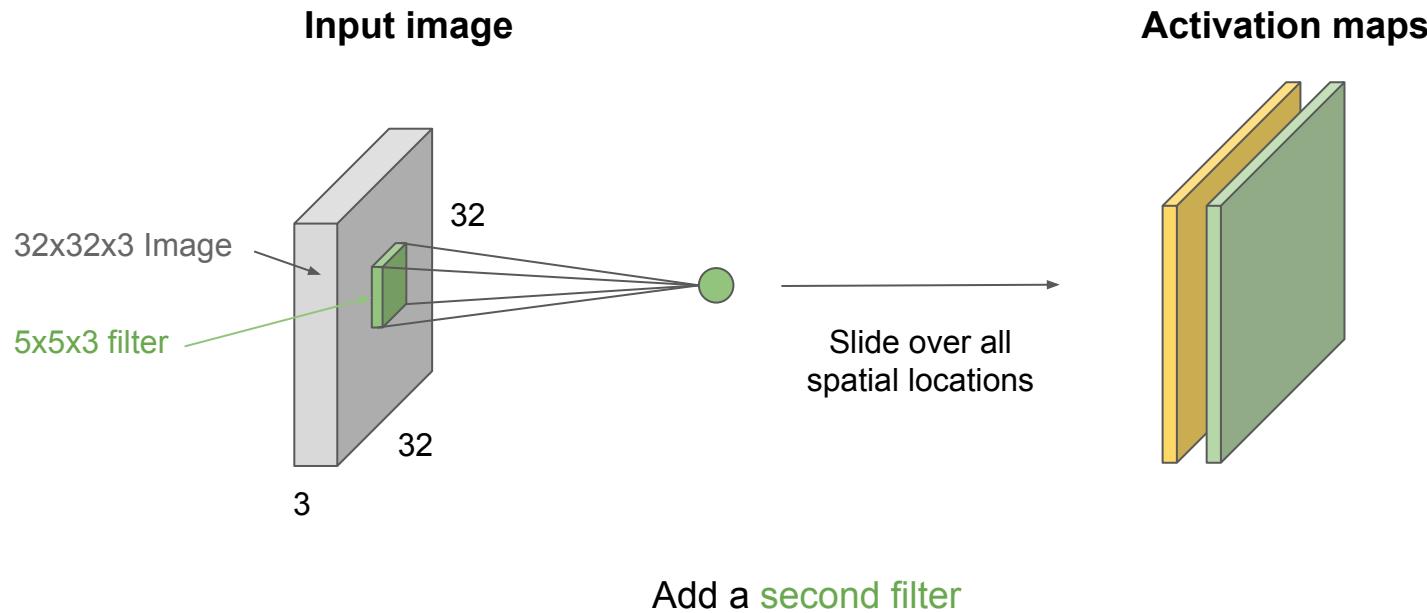
# Convolution layer



# Convolution layer

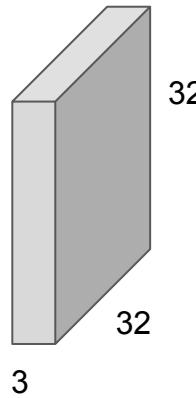


# Convolution layer



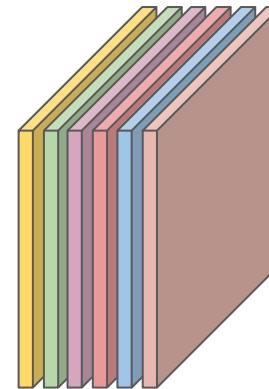
# Convolution layer

**Input image**

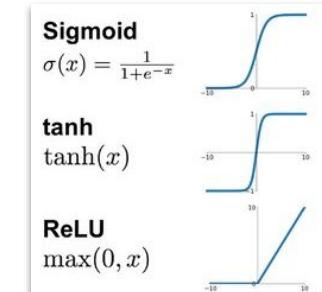
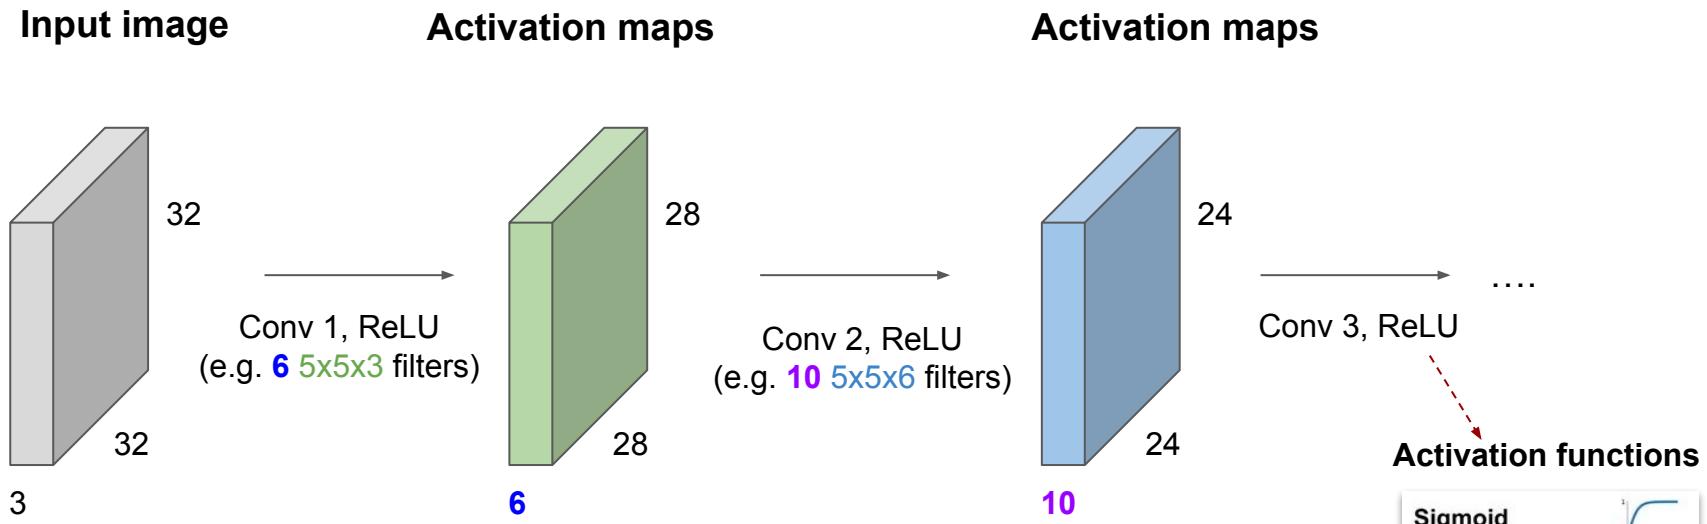


→  
A convolution layer with 6 filters

**Activation maps**



CNN is a sequence of convolution layers, interspersed with activation functions



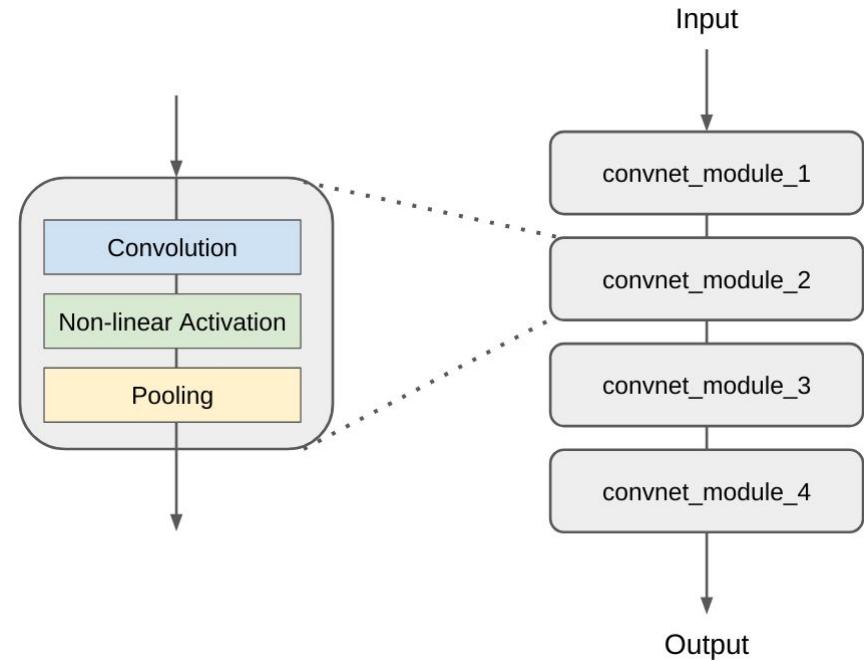
# Convolutional Neural Networks

A convolutional neural network (CNN) is a model  $f : \mathbb{R} \rightarrow \mathbb{R}$  that consists of a composition of  $L$  neural network layers that contain convolutions:

$$f(x) = f_L \circ f_{L-1} \circ \dots \circ f_1(x).$$

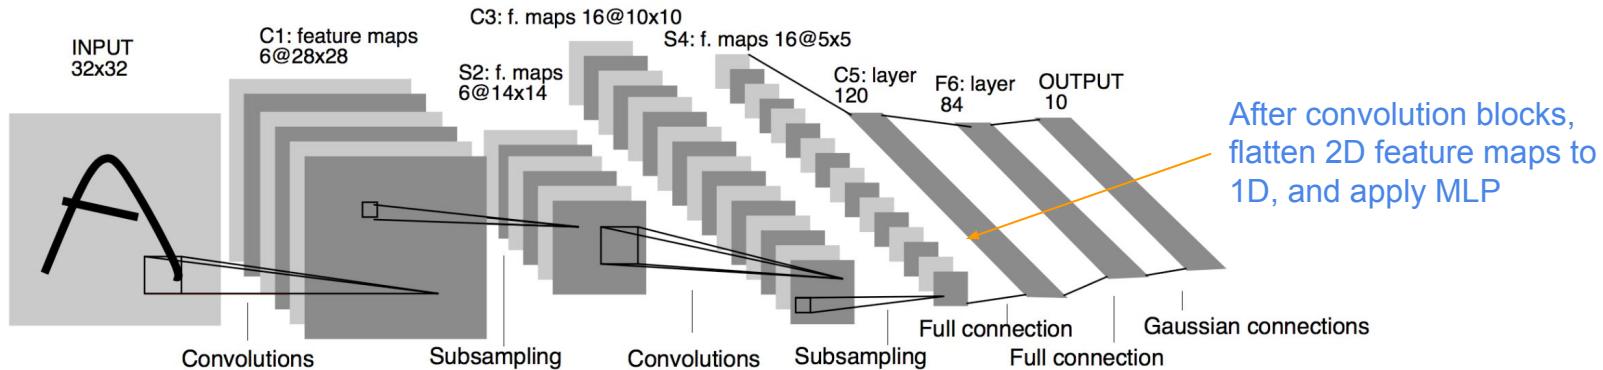
The final  $f_L$  is often a fully connected output layer of size one.

Typically, CNNs are made of consecutive convolution + activation + pooling layers that form into blocks.



# Example 1: LeNet for MNIST Digits Recognition

- LeNet successfully used CNNs for digit recognition [LeCun, Bottou, Bengio, Haffner 1998].



- LeNet contains multiple convolutional and max-pooling (called subsampling in the original paper) layers and a fully connected MLP in the end.
- Given an image of a handwritten digit, LeNet predicts it to be one of the 10 classes (0-9).
- This work inspired many subsequent CNN architectures.

# CNN in PyTorch

Layers definition

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

# CNN in PyTorch

Conv layer

Pooling layer

Fully-connecte  
d (linear) layer

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

# CNN in PyTorch

Function to call  
when running  
the layer

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

# CNN in PyTorch

Conv -> ReLU -> MaxPool

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

# CNN in PyTorch

The second conv block

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

# CNN in PyTorch

Flatten to 1D vectors

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

# CNN in PyTorch

Fully-connected neural network for classification

```
import torch.nn as nn
import torch.nn.functional as F

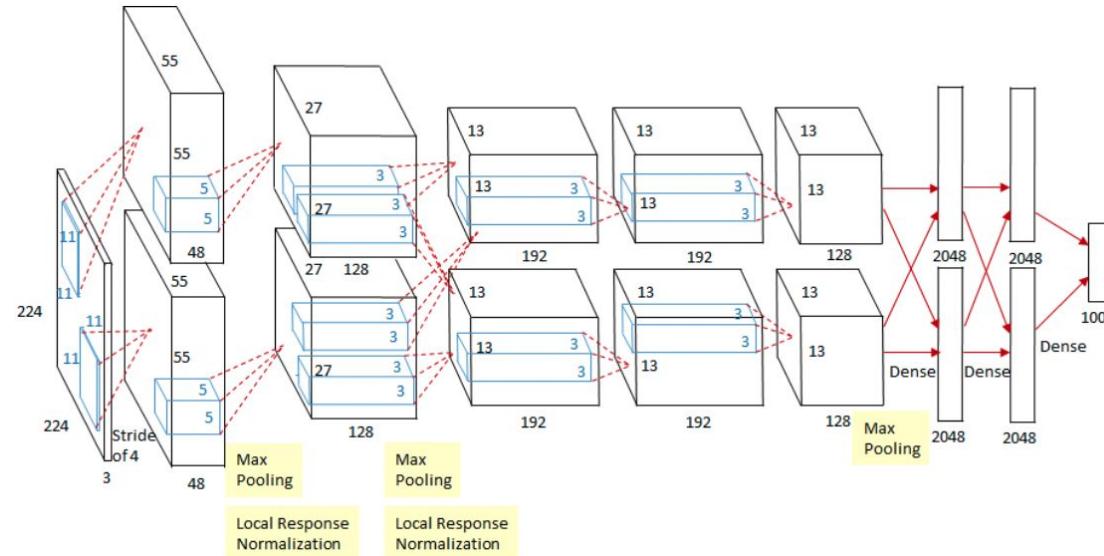
class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

# Example 2: AlexNet for ImageNet

- ImageNet classification with deep convolutional neural networks [Krizhevsky, Sutskever, Hinton 2012]



- AlexNet contains multiple convolutional layers followed by max pooling. The prediction is 1000 dimensional because that is the number of classes in the ImageNet classification task.
- Note that in the original implementation the image is separated into two parts that are in two different GPUs due to hardware limits.

# Example 3: ResNet

- Winner of the ImageNet Competition 2015
- Default architecture for many CNN models
- The “residual link” idea is widely used
- The most cited neural network of the 21st century

Data as of 2024/02/06

## Deep residual learning for image recognition

K He, X Zhang, S Ren, J Sun - ... and pattern recognition, 2016 - openaccess.thecvf.com |

... Deeper neural networks are more difficult to train. We present a residual learning frame to ease the training of networks that are substantially deeper than those used previously. ....

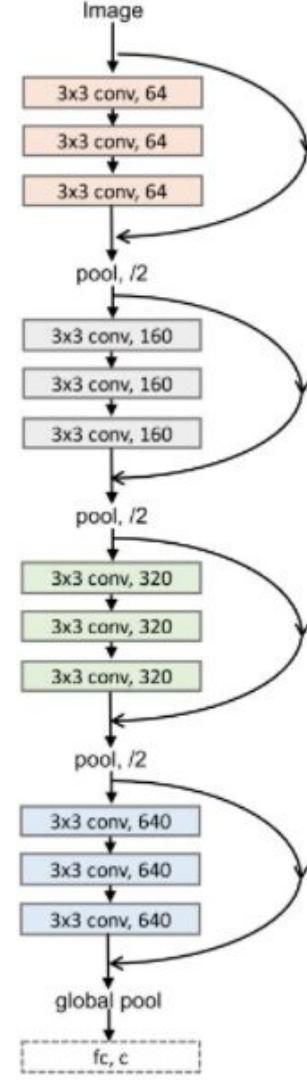
☆ Save 99 Cite Cited by 199144 Related articles All 76 versions Import into BibTeX

This lecture in 2023:

☆ Save 99 Cite Cited by 152947 Related articles All 73 versions

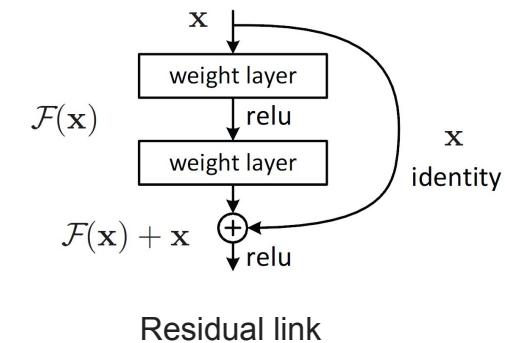
This lecture in 2022:

☆ Save 99 Cite Cited by 103017 Related articles All 63 versions



## Deep Residual Learning for Image Recognition

Kaiming He Xiangyu Zhang Shaoqing Ren Jian Sun  
Microsoft Research  
{kahe, v-xiangz, v-shren, jiansun}@microsoft.com



# What have been learned in CNN ? (for image classification)

- CNNs can be seen as extracting features at **low-, mid-, and high-levels** of abstraction that are relevant to corresponding visual concepts.
- Below, we reproduce **feature visualization** from the paper *Visualizing and Understanding Convolutional Networks* [Zeiler, Fergus 2013].

For a convolutional network with

Input → Conv Layer1 → Conv Layer2 → ⋯ → Conv Layer5 → ⋯ → Output,

we can visualize what each layer is doing by finding the image patch with highest activation responses.

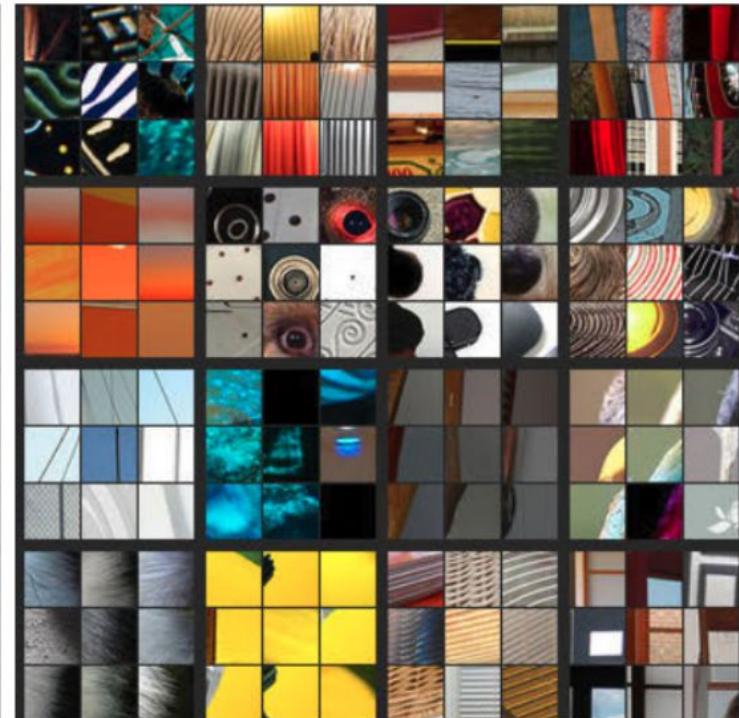
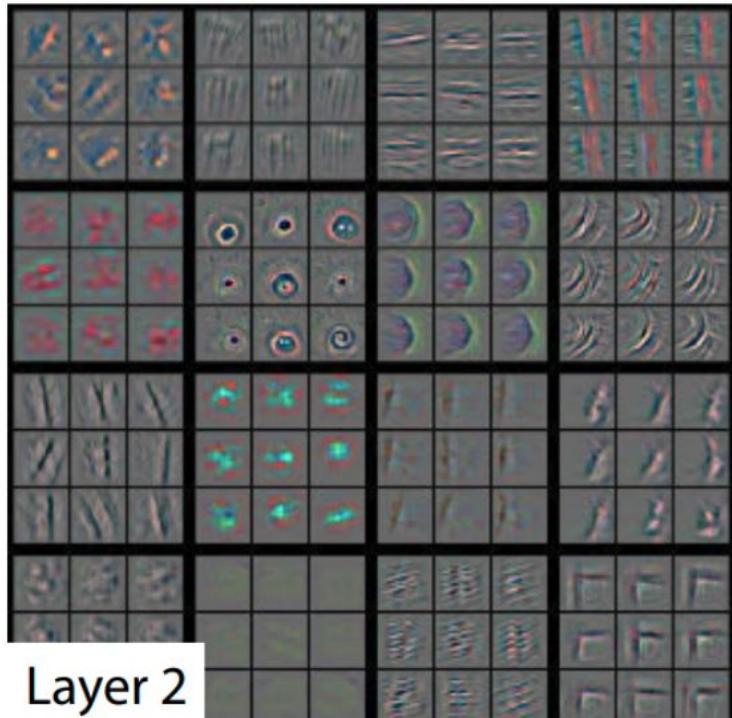
# What have been learned in CNN ? (for image classification)

- Layer 1 focuses on colored **edges** and **blobs**.



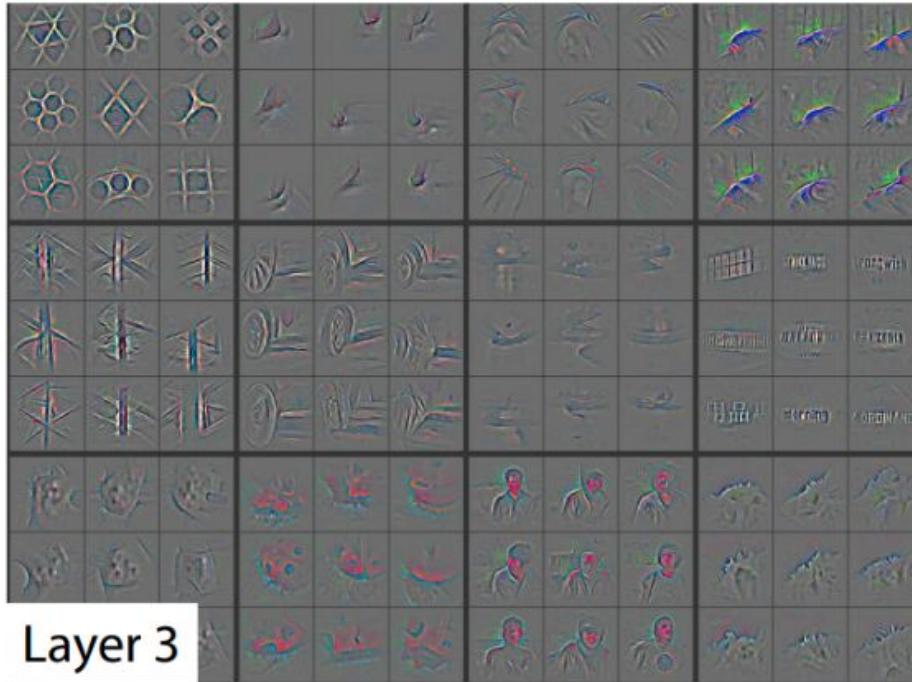
# What have been learned in CNN ? (for image classification)

- Layer 2-3 focus on **object parts**, such as the wheels of a car, or the beak of a bird.



# What have been learned in CNN ? (for image classification)

- Layer 2-3 focus on **object parts**, such as the wheels of a car, or the beak of a bird.

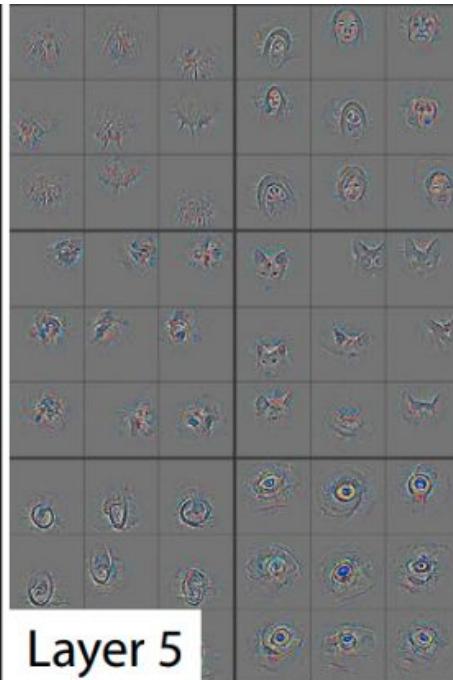


# What have been learned in CNN ? (for image classification)

- Layer 4-5 focus on **object parts** and even **entire objects**.



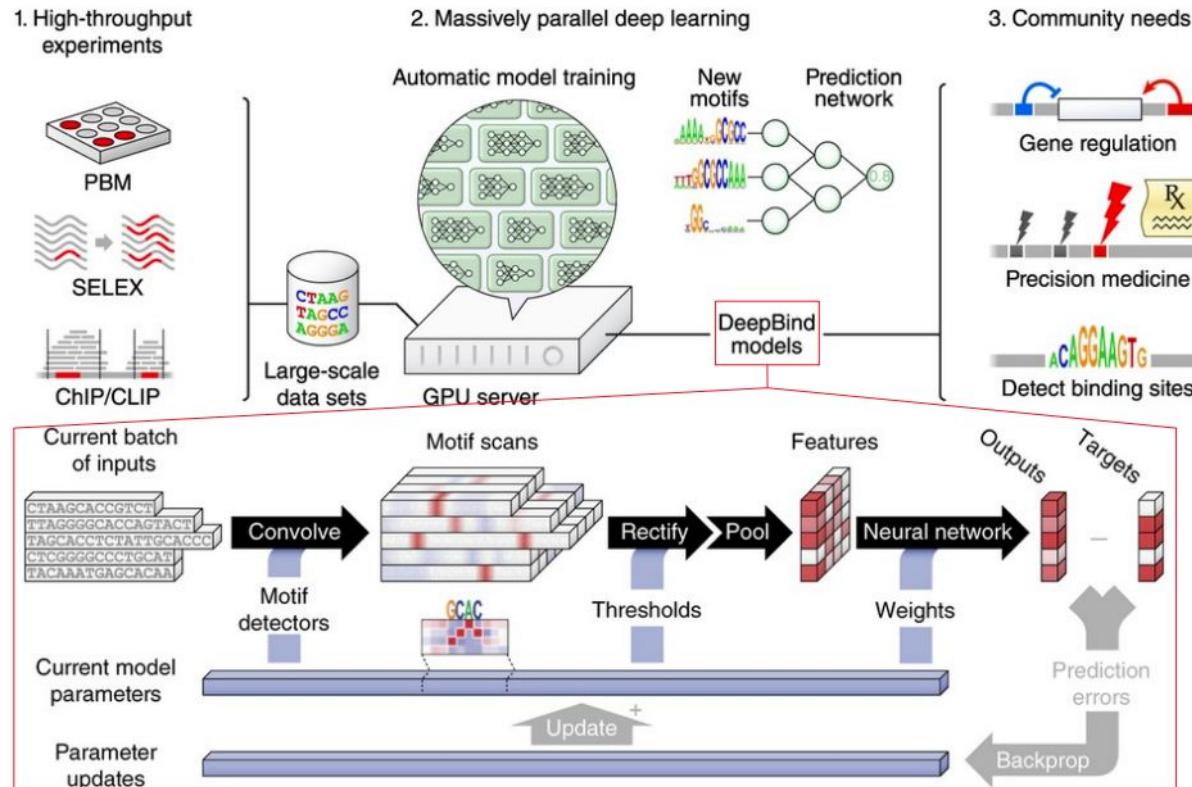
Layer 4



Layer 5

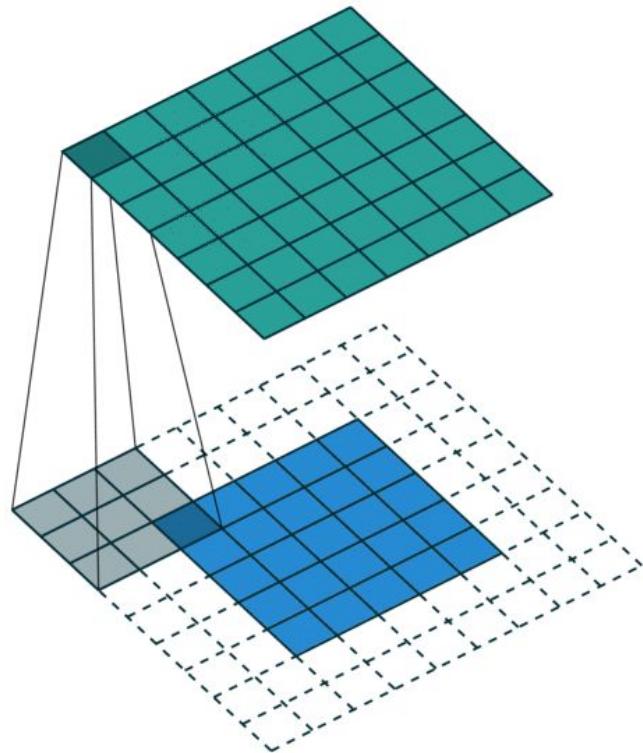


# Example: predicting DNA/RNA-protein binding



# Summary: key ideas of CNNs

- Use a set of filters to extract **local** features
- Use **multiple** filters to extract different features
- Spatially **share** parameters of each filter



# CNN summary

## CNNs

- Are used for all aspects of **computer vision**, and have won numerous pattern recognition competitions
- Able learn **interpretable features** at different levels of abstraction
- Typically, consist of **convolution layers**, **pooling layers**, **nonlinearities**, and **fully connected layers**

# A CSE Faculty Candidate Seminar Tomorrow (02/08)

Some of the coolest algorithms (e.g., sketching, graph theory, alignment) for biological sequence data analyses!

**Title:** Digital Genome Reconstruction with De Novo Assembly

**Speaker:** Haoyu Cheng (postdoctoral scholar at Dana-Farber Cancer Institute and Harvard Medical School)

**Location:** CODA Building, Second Floor, Room 230

**Date:** Thursday, February 8, 2024 at 11:00 am

**Abstract:**

<https://www.cse.gatech.edu/events/2024/02/08/cse-faculty-candidate-seminar-haoyu-cheng>

